

# Reinforcement Learning for Engineer-Mathematicians

A Comprehensive Guide to Theory and Applications

## About This Enhanced Edition

This comprehensive textbook bridges theory and practice in reinforcement learning:

- **18 Complete Chapters:** From mathematical foundations to research frontiers
- **Mathematical Rigor:** Formal theorems, proofs, and convergence analysis
- **Practical Focus:** Engineering applications and implementation guidance
- **Interactive Learning:** 5 Jupyter notebooks with Google Colab support
- **Modern Coverage:** Classical methods through deep reinforcement learning

Enhanced Edition 2024

Publisher Information

Copyright © 2024 Enhanced Edition

All rights reserved under Creative Commons Attribution-ShareAlike 4.0

**Open Educational Resource**

This work is licensed under CC BY-SA 4.0

**Repository:** <https://github.com/adiel2012/reinforcement-learning>

*To all engineers and mathematicians who seek to bridge  
the gap between theory and practice in the age of intelligent systems*

---

# Preface

This enhanced edition represents a comprehensive treatment of reinforcement learning designed specifically for engineer-mathematicians who require both theoretical rigor and practical implementation guidance. The field of reinforcement learning has evolved rapidly, with breakthroughs in deep learning enabling applications previously thought impossible.

## What Makes This Edition Special

This textbook bridges the gap between mathematical theory and engineering practice through:

- **Complete Coverage:** 18 chapters covering foundations through research frontiers
- **Mathematical Rigor:** Formal definitions, theorems, proofs, and convergence analysis
- **Engineering Focus:** Practical implementations and real-world applications
- **Interactive Learning:** Jupyter notebooks with Google Colab support
- **Modern Approach:** Classical methods through state-of-the-art deep RL

## How to Use This Book

The book is designed for flexible learning:

**Theory-First Learners** Begin with the mathematical foundations in Chapters 1-5, then explore advanced topics

**Hands-On Learners** Start with the interactive notebooks for Chapters 1-5, then delve into theory

**Practitioners** Focus on implementation chapters (6-8, 14-17) with theoretical background as needed

**Researchers** Use as a comprehensive reference for both classical and modern methods

---

## Acknowledgments

Special thanks to the reinforcement learning community for their open sharing of ideas, the developers of open-source tools that make interactive learning possible, and the countless researchers who have advanced this fascinating field.

# Contents

<b>Preface</b>	<b>v</b>
<b>I Mathematical Foundations</b>	<b>1</b>
<b>1 Introduction and Mathematical Prerequisites</b>	<b>5</b>
1.1 Motivation: From Control Theory to Learning Systems . . . . .	5
1.2 Mathematical Notation and Conventions . . . . .	6
1.2.1 Sets and Spaces . . . . .	7
1.2.2 Functions and Operators . . . . .	7
1.2.3 Probability and Expectation . . . . .	7
1.3 Probability Theory Refresher . . . . .	7
1.3.1 Probability Spaces and Random Variables . . . . .	8
1.3.2 Conditional Expectation and Martingales . . . . .	8
1.3.3 Concentration Inequalities . . . . .	8
1.4 Linear Algebra Essentials . . . . .	9
1.4.1 Vector Spaces and Norms . . . . .	9
1.4.2 Inner Products and Orthogonality . . . . .	10
1.4.3 Eigenvalues and Spectral Theory . . . . .	10
1.5 Optimization Fundamentals . . . . .	10
1.5.1 Convex Analysis . . . . .	11
1.5.2 Unconstrained Optimization . . . . .	11
1.5.3 Gradient Descent and Convergence Analysis . . . . .	11
1.6 Stochastic Processes and Markov Chains . . . . .	12
1.6.1 Discrete-Time Stochastic Processes . . . . .	12
1.6.2 Markov Chain Analysis . . . . .	12
1.6.3 Mixing Times and Convergence Rates . . . . .	12
1.7 Chapter Summary . . . . .	13
<b>2 Markov Decision Processes (MDPs)</b>	<b>15</b>
2.1 Understanding MDPs Through Examples . . . . .	15
2.2 Formal Definition and Mathematical Properties . . . . .	16

2.2.1	Assumptions and Regularity Conditions . . . . .	16
2.2.2	State and Action Spaces . . . . .	17
2.3	Policies and Value Functions . . . . .	17
2.3.1	Types of Policies . . . . .	17
2.3.2	Value Function Theory . . . . .	18
2.3.3	Bellman Equations . . . . .	19
2.4	Optimal Policies and Bellman Optimality . . . . .	19
2.4.1	Partial Ordering on Policies . . . . .	19
2.4.2	Bellman Optimality Equations . . . . .	19
2.5	Contraction Mapping Theorem and Fixed Points . . . . .	20
2.5.1	Bellman Operators . . . . .	20
2.5.2	Contraction Properties . . . . .	20
2.5.3	Banach Fixed Point Theorem Application . . . . .	20
2.6	Policy Improvement and Optimality . . . . .	21
2.6.1	Policy Improvement Theorem . . . . .	21
2.6.2	Policy Iteration Algorithm . . . . .	21
2.7	Computational Complexity Analysis . . . . .	21
2.7.1	Value Iteration Complexity . . . . .	21
2.7.2	Policy Iteration Complexity . . . . .	22
2.7.3	Modified Policy Iteration . . . . .	22
2.8	Connections to Classical Control Theory . . . . .	22
2.8.1	Linear Quadratic Regulator (LQR) . . . . .	22
2.8.2	Hamilton-Jacobi-Bellman Equation . . . . .	22
2.9	Advanced Topics . . . . .	23
2.9.1	Partially Observable MDPs (POMDPs) . . . . .	23
2.9.2	Constrained MDPs . . . . .	23
2.10	Chapter Summary . . . . .	23
<b>3</b>	<b>Dynamic Programming Foundations</b>	<b>25</b>
3.1	Principle of Optimality . . . . .	25
3.1.1	Mathematical Formulation . . . . .	25
3.1.2	Engineering Interpretation . . . . .	26
3.2	Value Iteration: Convergence Analysis . . . . .	26
3.2.1	Algorithm Description . . . . .	26
3.2.2	Convergence Theory . . . . .	26
3.2.3	Error Bounds and Stopping Criteria . . . . .	27
3.2.4	Computational Complexity . . . . .	27
3.3	Policy Iteration: Mathematical Guarantees . . . . .	28
3.3.1	Policy Evaluation . . . . .	28
3.3.2	Iterative Policy Evaluation . . . . .	28
3.3.3	Policy Improvement Analysis . . . . .	29



3.3.4	Global Convergence	29
3.4	Modified Policy Iteration	29
3.4.1	Algorithm and Convergence	29
3.4.2	Optimal Choice of Evaluation Steps	30
3.5	Asynchronous Dynamic Programming	30
3.5.1	Gauss-Seidel Value Iteration	30
3.5.2	Prioritized Sweeping	30
3.5.3	Real-Time Dynamic Programming	31
3.6	Linear Programming Formulation	31
3.6.1	Primal LP Formulation	31
3.6.2	Dual LP Formulation	32
3.7	Connections to Classical Control Theory	32
3.7.1	Discrete-Time Optimal Control	32
3.7.2	Stochastic Optimal Control	32
3.7.3	Model Predictive Control (MPC)	32
3.8	Computational Considerations	33
3.8.1	Curse of Dimensionality	33
3.8.2	Approximate Dynamic Programming	33
3.9	Chapter Summary	33
<b>4</b>	<b>Monte Carlo Methods</b>	<b>35</b>
4.1	Monte Carlo Estimation Theory	35
4.1.1	Basic Monte Carlo Principle	35
4.1.2	Application to Value Function Estimation	36
4.2	First-Visit vs. Every-Visit Methods	36
4.2.1	First-Visit Monte Carlo	36
4.2.2	Every-Visit Monte Carlo	36
4.3	Variance Reduction Techniques	37
4.3.1	Incremental Implementation	37
4.3.2	Baseline Subtraction	37
4.3.3	Control Variates	37
4.4	Importance Sampling in RL	37
4.4.1	Ordinary Importance Sampling	38
4.4.2	Weighted Importance Sampling	38
4.4.3	Per-Decision Importance Sampling	38
4.5	Off-Policy Monte Carlo Methods	38
4.5.1	Off-Policy Policy Evaluation	38
4.5.2	Off-Policy Monte Carlo Control	38
4.6	Convergence Analysis and Sample Complexity	40
4.6.1	Finite Sample Analysis	40
4.6.2	Asymptotic Convergence Rate	40

4.6.3	Sample Complexity . . . . .	40
4.7	Practical Considerations . . . . .	40
4.7.1	Exploration vs. Exploitation . . . . .	40
4.7.2	Function Approximation . . . . .	41
4.8	Chapter Summary . . . . .	41
<b>5</b>	<b>Temporal Difference Learning</b>	<b>43</b>
5.1	TD(0) Algorithm and Mathematical Analysis . . . . .	43
5.1.1	Basic TD(0) Update . . . . .	43
5.1.2	Relationship to Bellman Equation . . . . .	43
5.1.3	Convergence Analysis . . . . .	44
5.2	Bias-Variance Tradeoff in TD Methods . . . . .	44
5.2.1	Bias Analysis . . . . .	44
5.2.2	Variance Analysis . . . . .	45
5.2.3	Mean Squared Error Decomposition . . . . .	45
5.3	TD( $\lambda$ ) and Eligibility Traces . . . . .	45
5.3.1	Forward View: n-step Returns . . . . .	45
5.3.2	$\lambda$ -Return . . . . .	45
5.3.3	Backward View: Eligibility Traces . . . . .	46
5.3.4	Equivalence Theorem . . . . .	46
5.4	Convergence Theory for Linear Function Approximation . . . . .	47
5.4.1	Projected Bellman Equation . . . . .	47
5.4.2	Error Bounds . . . . .	47
5.5	Comparison with Monte Carlo and DP Methods . . . . .	47
5.5.1	Computational Complexity . . . . .	47
5.5.2	Sample Efficiency . . . . .	47
5.5.3	Bootstrapping vs. Sampling . . . . .	48
5.6	Advanced Topics . . . . .	48
5.6.1	Multi-step Methods . . . . .	48
5.6.2	True Online TD( $\lambda$ ) . . . . .	48
5.6.3	Gradient TD Methods . . . . .	48
5.7	Chapter Summary . . . . .	49
<b>II</b>	<b>Function Approximation</b>	<b>51</b>
<b>6</b>	<b>Q-Learning and SARSA Extensions</b>	<b>55</b>
6.1	Multi-Step Q-Learning . . . . .	55
6.1.1	n-Step Q-Learning . . . . .	55
6.1.2	Theoretical Analysis of n-Step Methods . . . . .	56
6.2	Q( $\lambda$ ) Learning . . . . .	57
6.2.1	Eligibility Traces for Q-Learning . . . . .	57

6.3	SARSA( $\lambda$ ) and True Online Methods . . . . .	57
6.3.1	SARSA( $\lambda$ ) Algorithm . . . . .	57
6.3.2	True Online SARSA( $\lambda$ ) . . . . .	57
6.4	Double Q-Learning . . . . .	58
6.4.1	The Maximization Bias Problem . . . . .	58
6.4.2	Double Q-Learning Algorithm . . . . .	58
6.4.3	Bias Reduction Analysis . . . . .	58
6.5	Expected SARSA . . . . .	59
6.5.1	Algorithm and Convergence . . . . .	59
6.6	Performance Analysis and Comparison . . . . .	59
6.6.1	Sample Complexity Bounds . . . . .	59
6.6.2	Empirical Comparison Framework . . . . .	60
6.7	Advanced Topics . . . . .	60
6.7.1	Gradient Q-Learning . . . . .	60
6.7.2	Distributional Q-Learning . . . . .	60
6.8	Implementation Considerations . . . . .	61
6.8.1	Memory and Computational Efficiency . . . . .	61
6.8.2	Hyperparameter Sensitivity . . . . .	61
6.9	Chapter Summary . . . . .	61
<b>7</b>	<b>Linear Function Approximation</b>	<b>63</b>
7.1	The Need for Generalization . . . . .	63
7.1.1	Limitations of Tabular Methods . . . . .	63
7.1.2	Function Approximation Framework . . . . .	64
7.2	Linear Function Approximation . . . . .	64
7.2.1	Feature Representation . . . . .	64
7.2.2	Linear Action-Value Approximation . . . . .	64
7.3	Gradient-Based Learning . . . . .	65
7.3.1	Value Function Approximation with Gradient Descent . . . . .	65
7.3.2	Stochastic Gradient Descent . . . . .	65
7.4	Temporal Difference with Function Approximation . . . . .	65
7.4.1	Semi-Gradient TD(0) . . . . .	65
7.4.2	Q-Learning with Function Approximation . . . . .	66
7.5	Convergence Analysis . . . . .	66
7.5.1	The Projection Matrix . . . . .	66
7.5.2	Fixed Point Analysis . . . . .	66
7.5.3	Mean Squared Error Bound . . . . .	67
7.6	The Deadly Triad . . . . .	67
7.6.1	Instability in Function Approximation . . . . .	67
7.6.2	Gradient-TD Methods . . . . .	67
7.7	Least-Squares Methods . . . . .	68

7.7.1	Least-Squares Temporal Difference (LSTD)	68
7.7.2	Recursive Least-Squares Implementation	68
7.8	Feature Construction	68
7.8.1	Tile Coding	68
7.8.2	Radial Basis Functions	69
7.8.3	Fourier Basis	69
7.9	Policy Gradient with Linear Approximation	69
7.9.1	Linear Policy Approximation	69
7.9.2	REINFORCE with Linear Features	70
7.10	Practical Considerations	70
7.10.1	Feature Scaling and Normalization	70
7.10.2	Regularization Techniques	70
7.11	Chapter Summary	70
<b>8</b>	<b>Neural Networks in Reinforcement Learning</b>	<b>73</b>
8.1	Neural Network Fundamentals for RL	73
8.1.1	Universal Function Approximation	73
8.1.2	Deep Neural Network Architecture	74
8.1.3	Gradient Computation via Backpropagation	74
8.2	Deep Q-Networks (DQN)	74
8.2.1	The DQN Algorithm	74
8.2.2	Experience Replay Analysis	74
8.2.3	Target Network Stabilization	75
8.3	Advanced DQN Variants	75
8.3.1	Double DQN	75
8.3.2	Dueling DQN	76
8.3.3	Prioritized Experience Replay	76
8.4	Policy-Based Deep RL	77
8.4.1	Deep Policy Networks	77
8.4.2	Deep Deterministic Policy Gradient (DDPG)	77
8.5	Challenges in Deep RL	78
8.5.1	Sample Efficiency	78
8.5.2	Stability and Convergence	78
8.5.3	Hyperparameter Sensitivity	78
8.6	Advanced Architectures	78
8.6.1	Recurrent Neural Networks for Partial Observability	78
8.6.2	Attention Mechanisms	79
8.6.3	Graph Neural Networks	79
8.7	Regularization and Generalization	79
8.7.1	Dropout in RL Networks	79
8.7.2	Batch Normalization	79

8.7.3	Weight Decay and L2 Regularization . . . . .	80
8.8	Theoretical Analysis . . . . .	80
8.8.1	Neural Tangent Kernel Theory . . . . .	80
8.8.2	Approximation Error Analysis . . . . .	80
8.9	Implementation Considerations . . . . .	80
8.9.1	Computational Graph and Automatic Differentiation . . . . .	80
8.9.2	Hardware Acceleration . . . . .	80
8.9.3	Memory Management . . . . .	81
8.10	Debugging and Diagnostics . . . . .	81
8.10.1	Common Deep RL Failures . . . . .	81
8.10.2	Ablation Studies . . . . .	81
8.11	Chapter Summary . . . . .	81
<b>III</b>	<b>Policy Methods</b>	<b>83</b>
<b>9</b>	<b>Policy Gradient Methods - Summary</b>	<b>87</b>
<b>10</b>	<b>Actor-Critic Methods - Summary</b>	<b>89</b>
<b>11</b>	<b>Advanced Policy Optimization - Summary</b>	<b>91</b>
<b>IV</b>	<b>Advanced Topics</b>	<b>93</b>
<b>12</b>	<b>Multi-Agent Reinforcement Learning - Summary</b>	<b>97</b>
<b>13</b>	<b>Hierarchical Reinforcement Learning - Summary</b>	<b>99</b>
<b>14</b>	<b>Model-Based Reinforcement Learning - Summary</b>	<b>101</b>
<b>15</b>	<b>Exploration and Exploitation - Summary</b>	<b>103</b>
<b>V</b>	<b>Applications and Frontiers</b>	<b>105</b>
<b>16</b>	<b>Transfer Learning and Meta-Learning - Summary</b>	<b>109</b>
<b>17</b>	<b>Real-World Applications and Deployment - Summary</b>	<b>111</b>
<b>18</b>	<b>Future Directions and Research Frontiers - Summary</b>	<b>113</b>
	<b>Conclusion</b>	<b>115</b>



## Part I

# Mathematical Foundations





---

This part establishes the mathematical foundations necessary for understanding reinforcement learning from both theoretical and engineering perspectives. We begin with essential mathematical prerequisites, develop the formal framework of Markov Decision Processes, and conclude with classical dynamic programming methods.

The treatment emphasizes mathematical rigor while maintaining practical relevance for engineering applications. Each concept is developed with careful attention to assumptions, proofs, and connections to control theory and optimization.



# Chapter 1

## Introduction and Mathematical Prerequisites

### Chapter Overview

This chapter introduces the fundamental mathematical tools needed for reinforcement learning and provides intuitive motivation for why RL represents a paradigm shift from classical control theory. We'll cover probability theory, linear algebra, optimization, and stochastic processes with practical examples.

### 1.1 Motivation: From Control Theory to Learning Systems

#### Why Reinforcement Learning?

Imagine teaching a child to ride a bicycle. You don't give them the equations of motion or tell them exactly how to balance. Instead, they learn through trial and error, gradually improving their balance and control. This is the essence of reinforcement learning.

Reinforcement learning represents a fundamental paradigm shift from classical control theory and optimization. While traditional engineering approaches rely on explicit models and well-defined objectives, reinforcement learning enables systems to learn optimal behavior through interaction with their environment.

#### Traditional Control vs. Reinforcement Learning

Consider a classic engineering problem: designing a controller for an inverted pendulum.

##### Traditional Control Approach:

1. Deriving the system dynamics using Lagrangian mechanics
2. Linearizing around the equilibrium point

3. Designing a feedback controller using pole placement or LQR
4. Implementing the controller with known parameters

**Reinforcement Learning Approach:**

1. Define states (angle, angular velocity) and actions (applied force)
2. Specify a reward function (positive for upright, negative for falling)
3. Allow the agent to explore and learn through trial and error
4. Converge to an optimal policy without explicit knowledge of dynamics

This fundamental difference opens up possibilities for systems where:

- Dynamics are unknown or too complex to model accurately
- Environment conditions change over time
- Multiple conflicting objectives must be balanced
- System parameters vary or degrade over time

**Industrial Example: Power Grid Management**

Modern power grids face unprecedented challenges with renewable energy integration, electric vehicle charging, and dynamic pricing. Traditional grid control relies on pre-computed lookup tables and heuristic rules. Reinforcement learning enables real-time optimization that adapts to:

- Variable renewable generation
- Changing demand patterns
- Equipment failures and network topology changes
- Market price fluctuations

## 1.2 Mathematical Notation and Conventions

**Notation Guide**

Throughout this book, we adopt consistent mathematical notation that aligns with both control theory and machine learning conventions. Don't worry if some symbols are unfamiliar now—we'll introduce them gradually with intuitive explanations.

Throughout this book, we adopt consistent mathematical notation that aligns with both control theory and machine learning conventions.

### 1.2.1 Sets and Spaces

#### Understanding Spaces

Think of a "space" as the collection of all possible values a variable can take. For example, if we're controlling a robot arm, the state space might include all possible joint angles and velocities.

$$\mathcal{S} = \text{State space (all possible states)} \quad (1.1)$$

$$\mathcal{A} = \text{Action space (all possible actions)} \quad (1.2)$$

$$\mathcal{R} = \text{Reward space (all possible rewards)} \quad (1.3)$$

$$\mathbb{R}^n = n\text{-dimensional real vector space} \quad (1.4)$$

$$\mathbb{R}^{m \times n} = \text{Space of } m \times n \text{ real matrices} \quad (1.5)$$

### 1.2.2 Functions and Operators

$$\pi : \mathcal{S} \rightarrow \mathcal{A} \quad (\text{Deterministic policy}) \quad (1.6)$$

$$\pi : \mathcal{S} \rightarrow \Delta(\mathcal{A}) \quad (\text{Stochastic policy}) \quad (1.7)$$

$$V^\pi : \mathcal{S} \rightarrow \mathbb{R} \quad (\text{Value function}) \quad (1.8)$$

$$Q^\pi : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R} \quad (\text{Action-value function}) \quad (1.9)$$

$$T : \mathbb{R}^{\mathcal{S}} \rightarrow \mathbb{R}^{\mathcal{S}} \quad (\text{Bellman operator}) \quad (1.10)$$

where  $\Delta(\mathcal{A})$  denotes the space of probability distributions over  $\mathcal{A}$ .

### 1.2.3 Probability and Expectation

$$\mathbb{P}(s'|s, a) = \text{Transition probability} \quad (1.11)$$

$$\mathbb{E}_\pi[\cdot] = \text{Expectation under policy } \pi \quad (1.12)$$

$$\mathbb{E}_{s \sim \mu}[\cdot] = \text{Expectation over distribution } \mu \quad (1.13)$$

## 1.3 Probability Theory Refresher

Reinforcement learning is fundamentally about making decisions under uncertainty. A solid understanding of probability theory is essential for analyzing convergence properties, sample complexity, and algorithm performance.

### 1.3.1 Probability Spaces and Random Variables

**Definition 1.1** (Probability Space). *A probability space is a triple  $(\Omega, \mathcal{F}, \mathbb{P})$  where:*

- $\Omega$  is the sample space (set of all possible outcomes)
- $\mathcal{F}$  is a  $\sigma$ -algebra on  $\Omega$  (collection of measurable events)
- $\mathbb{P} : \mathcal{F} \rightarrow [0, 1]$  is a probability measure satisfying:

1.  $\mathbb{P}(\Omega) = 1$

2. For disjoint events  $A_1, A_2, \dots$ :  $\mathbb{P}(\bigcup_{i=1}^{\infty} A_i) = \sum_{i=1}^{\infty} \mathbb{P}(A_i)$

**Definition 1.2** (Random Variable). *A random variable  $X$  is a measurable function  $X : \Omega \rightarrow \mathbb{R}$  such that for every Borel set  $B \subseteq \mathbb{R}$ , the preimage  $X^{-1}(B) \in \mathcal{F}$ .*

### 1.3.2 Conditional Expectation and Martingales

Conditional expectation plays a crucial role in reinforcement learning, particularly in the analysis of temporal difference methods and policy gradient algorithms.

**Definition 1.3** (Conditional Expectation). *Given random variables  $X$  and  $Y$ , the conditional expectation  $\mathbb{E}[X|Y]$  is the unique (almost surely) random variable that is:*

1. Measurable with respect to  $\sigma(Y)$
2. Satisfies  $\mathbb{E}[\mathbb{E}[X|Y] \cdot \mathbf{1}_A] = \mathbb{E}[X \cdot \mathbf{1}_A]$  for all  $A \in \sigma(Y)$

**Theorem 1.4** (Law of Total Expectation). *For random variables  $X$  and  $Y$ :*

$$\mathbb{E}[X] = \mathbb{E}[\mathbb{E}[X|Y]] \tag{1.14}$$

**Definition 1.5** (Martingale). *A sequence of random variables  $\{X_t\}_{t=0}^{\infty}$  is a martingale with respect to filtration  $\{\mathcal{F}_t\}_{t=0}^{\infty}$  if:*

1.  $X_t$  is  $\mathcal{F}_t$ -measurable for all  $t$
2.  $\mathbb{E}[|X_t|] < \infty$  for all  $t$
3.  $\mathbb{E}[X_{t+1}|\mathcal{F}_t] = X_t$  almost surely

Martingales are fundamental for proving convergence of stochastic algorithms in reinforcement learning.

### 1.3.3 Concentration Inequalities

Concentration inequalities provide bounds on the probability that random variables deviate from their expected values. These are essential for finite-sample analysis of RL algorithms.

**Theorem 1.6** (Hoeffding's Inequality). *Let  $X_1, \dots, X_n$  be independent random variables*

with  $X_i \in [a_i, b_i]$  almost surely. Then for any  $t > 0$ :

$$\mathbb{P} \left( \left| \frac{1}{n} \sum_{i=1}^n X_i - \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i] \right| \geq t \right) \leq 2 \exp \left( - \frac{2n^2 t^2}{\sum_{i=1}^n (b_i - a_i)^2} \right) \quad (1.15)$$

**Theorem 1.7** (Azuma's Inequality). *Let  $\{X_t\}_{t=0}^\infty$  be a martingale with respect to  $\{\mathcal{F}_t\}_{t=0}^\infty$  such that  $|X_{t+1} - X_t| \leq c_t$  almost surely. Then:*

$$\mathbb{P}(|X_n - X_0| \geq t) \leq 2 \exp \left( - \frac{t^2}{2 \sum_{i=0}^{n-1} c_i^2} \right) \quad (1.16)$$

## 1.4 Linear Algebra Essentials

Linear algebra provides the foundation for function approximation, policy parameterization, and many algorithmic techniques in reinforcement learning.

### 1.4.1 Vector Spaces and Norms

**Definition 1.8** (Vector Space). *A vector space  $V$  over field  $\mathbb{F}$  (typically  $\mathbb{R}$  or  $\mathbb{C}$ ) is a set equipped with vector addition and scalar multiplication satisfying:*

1. *Commutativity:*  $u + v = v + u$
2. *Associativity:*  $(u + v) + w = u + (v + w)$
3. *Identity:*  $\exists 0 \in V$  such that  $v + 0 = v$
4. *Inverse:*  $\forall v \in V, \exists -v$  such that  $v + (-v) = 0$
5. *Scalar associativity:*  $a(bv) = (ab)v$
6. *Scalar identity:*  $1v = v$
7. *Distributivity:*  $a(u + v) = au + av$  and  $(a + b)v = av + bv$

**Definition 1.9** (Norm). *A norm on vector space  $V$  is a function  $\|\cdot\| : V \rightarrow \mathbb{R}_{\geq 0}$  satisfying:*

1.  $\|v\| = 0$  if and only if  $v = 0$
2.  $\|av\| = |a|\|v\|$  for scalar  $a$
3.  $\|u + v\| \leq \|u\| + \|v\|$  (triangle inequality)

Common norms in  $\mathbb{R}^n$ :

$$\|x\|_1 = \sum_{i=1}^n |x_i| \quad (\ell_1 \text{ norm}) \quad (1.17)$$

$$\|x\|_2 = \sqrt{\sum_{i=1}^n x_i^2} \quad (\text{Euclidean norm}) \quad (1.18)$$

$$\|x\|_\infty = \max_{i=1,\dots,n} |x_i| \quad (\ell_\infty \text{ norm}) \quad (1.19)$$

### 1.4.2 Inner Products and Orthogonality

**Definition 1.10** (Inner Product). *An inner product on real vector space  $V$  is a function  $\langle \cdot, \cdot \rangle : V \times V \rightarrow \mathbb{R}$  satisfying:*

1. *Symmetry:*  $\langle u, v \rangle = \langle v, u \rangle$
2. *Linearity:*  $\langle au + bv, w \rangle = a\langle u, w \rangle + b\langle v, w \rangle$
3. *Positive definiteness:*  $\langle v, v \rangle \geq 0$  with equality iff  $v = 0$

The induced norm is  $\|v\| = \sqrt{\langle v, v \rangle}$ .

**Theorem 1.11** (Cauchy-Schwarz Inequality). *For vectors  $u, v$  in an inner product space:*

$$|\langle u, v \rangle| \leq \|u\| \|v\| \quad (1.20)$$

*with equality if and only if  $u$  and  $v$  are linearly dependent.*

### 1.4.3 Eigenvalues and Spectral Theory

**Definition 1.12** (Eigenvalue and Eigenvector). *For matrix  $A \in \mathbb{R}^{n \times n}$ , scalar  $\lambda$  is an eigenvalue with corresponding eigenvector  $v \neq 0$  if:*

$$Av = \lambda v \quad (1.21)$$

**Theorem 1.13** (Spectral Theorem for Symmetric Matrices). *Every real symmetric matrix  $A$  has an orthonormal basis of eigenvectors with real eigenvalues. If  $A = Q\Lambda Q^T$  where  $Q$  is orthogonal and  $\Lambda$  is diagonal, then:*

$$A = \sum_{i=1}^n \lambda_i q_i q_i^T \quad (1.22)$$

*where  $\lambda_i$  are eigenvalues and  $q_i$  are corresponding orthonormal eigenvectors.*

## 1.5 Optimization Fundamentals

Optimization theory underpins virtually all reinforcement learning algorithms, from value iteration to policy gradient methods.



### 1.5.1 Convex Analysis

**Definition 1.14** (Convex Set). A set  $C \subseteq \mathbb{R}^n$  is convex if for all  $x, y \in C$  and  $\lambda \in [0, 1]$ :

$$\lambda x + (1 - \lambda)y \in C \quad (1.23)$$

**Definition 1.15** (Convex Function). A function  $f : \mathbb{R}^n \rightarrow \mathbb{R}$  is convex if its domain is convex and for all  $x, y$  in the domain and  $\lambda \in [0, 1]$ :

$$f(\lambda x + (1 - \lambda)y) \leq \lambda f(x) + (1 - \lambda)f(y) \quad (1.24)$$

**Theorem 1.16** (First-Order Characterization of Convexity). For differentiable function  $f$ , the following are equivalent:

1.  $f$  is convex
2.  $f(y) \geq f(x) + \nabla f(x)^T(y - x)$  for all  $x, y$
3.  $\nabla f$  is monotone:  $(\nabla f(x) - \nabla f(y))^T(x - y) \geq 0$

### 1.5.2 Unconstrained Optimization

**Theorem 1.17** (Necessary Conditions for Optimality). If  $x^*$  is a local minimum of differentiable function  $f$ , then:

$$\nabla f(x^*) = 0 \quad (1.25)$$

If  $f$  is twice differentiable, then additionally:

$$\nabla^2 f(x^*) \succeq 0 \quad (\text{positive semidefinite}) \quad (1.26)$$

**Theorem 1.18** (Sufficient Conditions for Optimality). If  $\nabla f(x^*) = 0$  and  $\nabla^2 f(x^*) \succ 0$  (positive definite), then  $x^*$  is a strict local minimum.

### 1.5.3 Gradient Descent and Convergence Analysis

The gradient descent algorithm iterates:

$$x_{k+1} = x_k - \alpha_k \nabla f(x_k) \quad (1.27)$$

**Theorem 1.19** (Convergence of Gradient Descent). For convex function  $f$  with  $L$ -Lipschitz gradient and step size  $\alpha \leq 1/L$ :

$$f(x_k) - f(x^*) \leq \frac{\|x_0 - x^*\|^2}{2\alpha k} \quad (1.28)$$

where  $x^*$  is the optimal solution.

For strongly convex functions, the convergence rate improves to exponential.

## 1.6 Stochastic Processes and Markov Chains

Understanding stochastic processes is crucial for analyzing the temporal dynamics of reinforcement learning systems.

### 1.6.1 Discrete-Time Stochastic Processes

**Definition 1.20** (Stochastic Process). *A discrete-time stochastic process is a sequence of random variables  $\{X_t\}_{t=0}^{\infty}$  where each  $X_t$  takes values in some state space  $\mathcal{S}$ .*

**Definition 1.21** (Markov Property). *A stochastic process  $\{X_t\}_{t=0}^{\infty}$  satisfies the Markov property if:*

$$\mathbb{P}(X_{t+1} = s' | X_t = s, X_{t-1} = s_{t-1}, \dots, X_0 = s_0) = \mathbb{P}(X_{t+1} = s' | X_t = s) \quad (1.29)$$

for all states  $s, s', s_0, \dots, s_{t-1}$  and times  $t \geq 0$ .

### 1.6.2 Markov Chain Analysis

For finite state space  $\mathcal{S} = \{1, 2, \dots, n\}$ , a Markov chain is characterized by its transition matrix  $P \in \mathbb{R}^{n \times n}$  where  $P_{ij} = \mathbb{P}(X_{t+1} = j | X_t = i)$ .

**Definition 1.22** (Irreducibility and Aperiodicity). *A Markov chain is:*

- **Irreducible** if every state is reachable from every other state
- **Aperiodic** if  $\gcd\{n \geq 1 : P_{ii}^{(n)} > 0\} = 1$  for some state  $i$

**Theorem 1.23** (Fundamental Theorem of Markov Chains). *For an irreducible, aperiodic, finite Markov chain:*

1. *There exists a unique stationary distribution  $\pi$  satisfying  $\pi = \pi P$*
2.  *$\lim_{t \rightarrow \infty} P^t = \mathbf{1}\pi^T$  where  $\mathbf{1}$  is the vector of ones*
3. *For any initial distribution  $\mu_0$ :  $\lim_{t \rightarrow \infty} \|\mu_t - \pi\|_{TV} = 0$*

### 1.6.3 Mixing Times and Convergence Rates

**Definition 1.24** (Total Variation Distance). *The total variation distance between distributions  $\mu$  and  $\nu$  on finite space  $\mathcal{S}$  is:*

$$\|\mu - \nu\|_{TV} = \frac{1}{2} \sum_{s \in \mathcal{S}} |\mu(s) - \nu(s)| \quad (1.30)$$

**Definition 1.25** (Mixing Time). *The mixing time of a Markov chain is:*

$$t_{mix}(\epsilon) = \min\{t : \max_{i \in \mathcal{S}} \|P^t(i, \cdot) - \pi\|_{TV} \leq \epsilon\} \quad (1.31)$$

Understanding mixing times is essential for analyzing sample complexity in reinforcement

learning algorithms that rely on sampling from stationary distributions.

## 1.7 Chapter Summary

This chapter established the mathematical foundations necessary for rigorous analysis of reinforcement learning algorithms. Key concepts include:

- The paradigm shift from model-based control to learning-based optimization
- Probability theory tools: conditional expectation, martingales, concentration inequalities
- Linear algebra foundations: vector spaces, norms, spectral theory
- Convex optimization and gradient descent convergence analysis
- Markov chain theory and convergence to stationary distributions

These mathematical tools will be applied throughout the book to analyze algorithm convergence, sample complexity, and performance guarantees. The next chapter develops the formal framework of Markov Decision Processes, which provides the mathematical foundation for all subsequent reinforcement learning algorithms.



## Chapter 2

# Markov Decision Processes (MDPs)

### Chapter Overview

This chapter introduces Markov Decision Processes (MDPs), the mathematical foundation of reinforcement learning. We'll build intuition through concrete examples before diving into the formal theory, then explore solution methods like value iteration and policy iteration.

### What is an MDP?

Think of an MDP as a mathematical description of a decision-making situation where:

- You observe the current situation (state)
- You choose an action based on what you observe
- The world responds by transitioning to a new state and giving you a reward
- This process repeats over time

The key insight is that the future only depends on the current state, not the entire history—this is the Markov property.

Markov Decision Processes provide the mathematical framework for modeling sequential decision-making under uncertainty. This chapter develops the formal theory of MDPs with particular attention to mathematical rigor and engineering applications.

## 2.1 Understanding MDPs Through Examples

Before diving into formal definitions, let's build intuition through a concrete example.

### Grid World Navigation

Consider a robot navigating a  $4 \times 4$  grid world:

- **States:** Each cell in the grid (16 total states)
- **Actions:** Move up, down, left, or right
- **Transitions:** Move to adjacent cell (or stay put if hitting a wall)
- **Rewards:** +10 for reaching the goal, -1 for each step, -10 for falling into holes
- **Goal:** Find the shortest path to the target while avoiding obstacles

## 2.2 Formal Definition and Mathematical Properties

Now that we have intuition, let's formalize these concepts.

**Definition 2.1** (Markov Decision Process). *A Markov Decision Process is a 5-tuple  $(\mathcal{S}, \mathcal{A}, P, \mathcal{R}, \gamma)$  where:*

- $\mathcal{S}$  is the **state space** (all possible situations)
- $\mathcal{A}$  is the **action space** (all possible actions)
- $P : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$  is the **transition kernel** (dynamics)
- $\mathcal{R} : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  is the **reward function** (immediate feedback)
- $\gamma \in [0, 1)$  is the **discount factor** (how much we value future rewards)

### Understanding the Components

- **State space  $\mathcal{S}$ :** All possible configurations of your system
- **Action space  $\mathcal{A}$ :** All decisions you can make in any given state
- **Transition function  $P$ :** Describes how actions change states (the "physics" of your world)
- **Reward function  $\mathcal{R}$ :** Immediate feedback telling you how good an action was
- **Discount factor  $\gamma$ :** How much you care about future vs. immediate rewards (0 = only care about immediate, close to 1 = care about long-term)

The transition kernel satisfies  $\sum_{s' \in \mathcal{S}} P(s, a, s') = 1$  for all  $(s, a) \in \mathcal{S} \times \mathcal{A}$ , and we write  $P(s'|s, a) = P(s, a, s')$  for the probability of transitioning to state  $s'$  from state  $s$  under action  $a$ .

### 2.2.1 Assumptions and Regularity Conditions

For mathematical tractability, we typically assume:

**Assumption 2.2** (Measurability). *The state and action spaces are measurable spaces, and the transition kernel and reward function are measurable with respect to the appropriate  $\sigma$ -algebras.*

**Assumption 2.3** (Bounded Rewards). *The reward function satisfies  $\sup_{s,a} |\mathcal{R}(s,a)| \leq R_{\max} < \infty$ .*

**Assumption 2.4** (Discount Factor). *The discount factor satisfies  $\gamma \in [0, 1)$  to ensure convergence of infinite-horizon value functions.*

### 2.2.2 State and Action Spaces

#### Discrete Spaces

For finite MDPs with  $|\mathcal{S}| = n$  and  $|\mathcal{A}| = m$ , we can represent:

- Transition probabilities as tensors  $P^a \in \mathbb{R}^{n \times n}$  for each action  $a$
- Rewards as matrices  $R \in \mathbb{R}^{n \times m}$
- Policies as matrices  $\Pi \in [0, 1]^{n \times m}$  with  $\sum_a \Pi(s, a) = 1$

#### Continuous Spaces

For continuous state spaces  $\mathcal{S} \subseteq \mathbb{R}^d$ , the transition kernel becomes a probability measure:

$$P(\cdot | s, a) : \mathcal{B}(\mathcal{S}) \rightarrow [0, 1] \quad (2.1)$$

where  $\mathcal{B}(\mathcal{S})$  is the Borel  $\sigma$ -algebra on  $\mathcal{S}$ .

#### Engineering Example: Inverted Pendulum

Consider an inverted pendulum with:

- State:  $s = (\theta, \dot{\theta}) \in [-\pi, \pi] \times [-10, 10]$  (angle and angular velocity)
- Action:  $a \in [-5, 5]$  (applied torque)
- Dynamics:  $\ddot{\theta} = \frac{g}{l} \sin \theta + \frac{a}{ml^2}$  plus noise
- Reward:  $r(s, a) = -\theta^2 - 0.1\dot{\theta}^2 - 0.01a^2$  (quadratic cost)

## 2.3 Policies and Value Functions

### What is a Policy?

A policy is simply a decision-making rule. It tells an agent what action to take in each possible state. Think of it as a strategy or game plan.

### 2.3.1 Types of Policies

**Definition 2.5** (Deterministic Policy). *A deterministic policy is a function  $\pi : \mathcal{S} \rightarrow \mathcal{A}$  that maps each state to exactly one action.*

### Deterministic Policy Example

In our grid world: "Always move towards the goal" could be a deterministic policy where  $\pi(\text{state}) = \text{direction\_to\_goal}$ .

**Definition 2.6** (Stochastic Policy). *A stochastic policy  $\pi : \mathcal{S} \rightarrow \Delta(\mathcal{A})$  assigns a probability distribution over actions for each state, where  $\Delta(\mathcal{A})$  is the space of probability measures on  $\mathcal{A}$ .*

### Stochastic Policy Example

In grid world: "Move towards goal with probability 0.8, move randomly otherwise" gives  $\pi(\text{best\_action}|s) = 0.8$  and equal probability to other actions.

**Definition 2.7** (History-Dependent Policy). *A history-dependent policy depends on the entire sequence of past states and actions:  $\pi_t : (\mathcal{S} \times \mathcal{A})^t \times \mathcal{S} \rightarrow \Delta(\mathcal{A})$*

### Why Focus on Markovian Policies?

While policies could potentially use the entire history, the Markov property means that optimal policies only need to depend on the current state. This greatly simplifies our analysis!

**Theorem 2.8** (Sufficiency of Markovian Policies). *For any history-dependent policy, there exists a Markovian policy that achieves the same expected discounted reward.*

*Proof.* This follows from the Markov property of the state transitions. The expected future reward depends only on the current state, not the history of how that state was reached.  $\square$

## 2.3.2 Value Function Theory

**Definition 2.9** (State Value Function). *The state value function for policy  $\pi$  is:*

$$V^\pi(s) = \mathbb{E}^\pi \left[ \sum_{t=0}^{\infty} \gamma^t \mathcal{R}(S_t, A_t) \mid S_0 = s \right] \quad (2.2)$$

**Definition 2.10** (Action Value Function). *The action value function ( $Q$ -function) for policy  $\pi$  is:*

$$Q^\pi(s, a) = \mathbb{E}^\pi \left[ \sum_{t=0}^{\infty} \gamma^t \mathcal{R}(S_t, A_t) \mid S_0 = s, A_0 = a \right] \quad (2.3)$$

**Theorem 2.11** (Existence and Uniqueness of Value Functions). *Under Assumptions 1-3, the value functions  $V^\pi$  and  $Q^\pi$  exist, are unique, and satisfy  $\|V^\pi\|_\infty \leq \frac{R_{max}}{1-\gamma}$ .*

*Proof.* The geometric series  $\sum_{t=0}^{\infty} \gamma^t R_{max}$  converges to  $\frac{R_{max}}{1-\gamma}$  since  $\gamma < 1$ . Uniqueness follows from the linearity of expectation.  $\square$



### 2.3.3 Bellman Equations

The fundamental recursive relationships in reinforcement learning are the Bellman equations.

**Theorem 2.12** (Bellman Equations for Policy Evaluation). *For any policy  $\pi$ :*

$$V^\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left[ \mathcal{R}(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^\pi(s') \right] \quad (2.4)$$

$$Q^\pi(s, a) = \mathcal{R}(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) \sum_{a' \in \mathcal{A}} \pi(a'|s') Q^\pi(s', a') \quad (2.5)$$

*Proof.* By the tower rule of conditional expectation:

$$V^\pi(s) = \mathbb{E}^\pi \left[ \mathcal{R}(S_0, A_0) + \gamma \sum_{t=1}^{\infty} \gamma^{t-1} \mathcal{R}(S_t, A_t) \mid S_0 = s \right] \quad (2.6)$$

$$= \mathbb{E}^\pi[\mathcal{R}(S_0, A_0) \mid S_0 = s] + \gamma \mathbb{E}^\pi[V^\pi(S_1) \mid S_0 = s] \quad (2.7)$$

Expanding the expectations gives the Bellman equation.  $\square$

## 2.4 Optimal Policies and Bellman Optimality

### 2.4.1 Partial Ordering on Policies

**Definition 2.13** (Policy Partial Order). *Policy  $\pi_1$  dominates policy  $\pi_2$  (written  $\pi_1 \geq \pi_2$ ) if:*

$$V^{\pi_1}(s) \geq V^{\pi_2}(s) \quad \forall s \in \mathcal{S} \quad (2.8)$$

**Theorem 2.14** (Existence of Optimal Policies). *There exists an optimal deterministic policy  $\pi^*$  such that:*

$$V^{\pi^*}(s) = \max_{\pi} V^{\pi}(s) \equiv V^*(s) \quad \forall s \in \mathcal{S} \quad (2.9)$$

### 2.4.2 Bellman Optimality Equations

**Theorem 2.15** (Bellman Optimality Equations). *The optimal value functions satisfy:*

$$V^*(s) = \max_{a \in \mathcal{A}} \left[ \mathcal{R}(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^*(s') \right] \quad (2.10)$$

$$Q^*(s, a) = \mathcal{R}(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) \max_{a' \in \mathcal{A}} Q^*(s', a') \quad (2.11)$$

**Corollary 2.16** (Optimal Policy Extraction). *An optimal policy can be extracted from the optimal value functions as:*

$$\pi^*(s) \in \operatorname{argmax}_{a \in \mathcal{A}} Q^*(s, a) \quad (2.12)$$

## 2.5 Contraction Mapping Theorem and Fixed Points

The mathematical foundation for proving convergence of dynamic programming algorithms relies on contraction mapping theory.

### 2.5.1 Bellman Operators

**Definition 2.17** (Bellman Operator). *For policy  $\pi$ , the Bellman operator  $T^\pi : \mathbb{R}^{\mathcal{S}} \rightarrow \mathbb{R}^{\mathcal{S}}$  is defined by:*

$$(T^\pi V)(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left[ \mathcal{R}(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s') \right] \quad (2.13)$$

**Definition 2.18** (Bellman Optimality Operator). *The Bellman optimality operator  $T^* : \mathbb{R}^{\mathcal{S}} \rightarrow \mathbb{R}^{\mathcal{S}}$  is defined by:*

$$(T^* V)(s) = \max_{a \in \mathcal{A}} \left[ \mathcal{R}(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s') \right] \quad (2.14)$$

### 2.5.2 Contraction Properties

**Theorem 2.19** (Contraction Property of Bellman Operators). *Under the supremum norm  $\|V\|_\infty = \max_{s \in \mathcal{S}} |V(s)|$ :*

1.  $T^\pi$  is a  $\gamma$ -contraction:  $\|T^\pi V_1 - T^\pi V_2\|_\infty \leq \gamma \|V_1 - V_2\|_\infty$
2.  $T^*$  is a  $\gamma$ -contraction:  $\|T^* V_1 - T^* V_2\|_\infty \leq \gamma \|V_1 - V_2\|_\infty$

*Proof.* For the policy operator:

$$|(T^\pi V_1)(s) - (T^\pi V_2)(s)| = \left| \sum_a \pi(a|s) \gamma \sum_{s'} P(s'|s, a) [V_1(s') - V_2(s')] \right| \quad (2.15)$$

$$\leq \sum_a \pi(a|s) \gamma \sum_{s'} P(s'|s, a) |V_1(s') - V_2(s')| \quad (2.16)$$

$$\leq \gamma \|V_1 - V_2\|_\infty \sum_a \pi(a|s) \sum_{s'} P(s'|s, a) \quad (2.17)$$

$$= \gamma \|V_1 - V_2\|_\infty \quad (2.18)$$

The proof for  $T^*$  follows similarly using the fact that the max operator is non-expansive.  $\square$

### 2.5.3 Banach Fixed Point Theorem Application

**Theorem 2.20** (Banach Fixed Point Theorem). *Let  $(X, d)$  be a complete metric space and  $T : X \rightarrow X$  be a contraction mapping with contraction factor  $\gamma < 1$ . Then:*

1.  $T$  has a unique fixed point  $x^* \in X$

2. For any  $x_0 \in X$ , the sequence  $x_{n+1} = T(x_n)$  converges to  $x^*$

3. The convergence rate is geometric:  $d(x_n, x^*) \leq \gamma^n d(x_0, x^*)$

**Corollary 2.21** (Convergence of Value Iteration). *The value iteration algorithm  $V_{k+1} = T^*V_k$  converges geometrically to the unique optimal value function  $V^*$  at rate  $\gamma$ .*

## 2.6 Policy Improvement and Optimality

### 2.6.1 Policy Improvement Theorem

**Theorem 2.22** (Policy Improvement Theorem). *Let  $\pi$  be any policy and define the improved policy  $\pi'$  by:*

$$\pi'(s) \in \operatorname{argmax}_{a \in \mathcal{A}} Q^\pi(s, a) \quad (2.19)$$

*Then  $\pi' \geq \pi$ , with strict inequality unless  $\pi$  is optimal.*

*Proof.* For any state  $s$ :

$$Q^\pi(s, \pi'(s)) \geq Q^\pi(s, \pi(s)) = V^\pi(s) \quad (2.20)$$

By the policy evaluation equation and induction, this implies  $V^{\pi'}(s) \geq V^\pi(s)$ .  $\square$

### 2.6.2 Policy Iteration Algorithm

The policy improvement theorem leads to the policy iteration algorithm:

---

**Algorithm 1** Policy Iteration

---

**Require:** Initial policy  $\pi_0$

**Ensure:** Optimal policy  $\pi^*$

$k \leftarrow 0$

**repeat**

**Policy Evaluation:** Solve  $V^{\pi_k} = T^{\pi_k} V^{\pi_k}$

**Policy Improvement:**  $\pi_{k+1}(s) \leftarrow \operatorname{argmax}_a Q^{\pi_k}(s, a)$

$k \leftarrow k + 1$

**until**  $\pi_k = \pi_{k-1}$

**return**  $\pi^* = \pi_k$

---

**Theorem 2.23** (Convergence of Policy Iteration). *Policy iteration converges to an optimal policy in finitely many iterations for finite MDPs.*

## 2.7 Computational Complexity Analysis

### 2.7.1 Value Iteration Complexity

For finite MDPs with  $|\mathcal{S}| = n$  and  $|\mathcal{A}| = m$ :

- **Time per iteration:**  $O(mn^2)$  operations

- **Iterations to  $\epsilon$ -accuracy:**  $O(\log(\epsilon^{-1}))$  iterations
- **Total complexity:**  $O(mn^2 \log(\epsilon^{-1}))$

### 2.7.2 Policy Iteration Complexity

- **Policy evaluation:**  $O(n^3)$  for direct matrix inversion,  $O(n^2)$  per iteration for iterative methods
- **Policy improvement:**  $O(mn^2)$
- **Number of policy iterations:** At most  $m^n$  (typically much smaller)

### 2.7.3 Modified Policy Iteration

To balance the computational costs, modified policy iteration performs only  $k$  steps of policy evaluation:

---

**Algorithm 2** Modified Policy Iteration
 

---

**Require:** Initial policy  $\pi_0$ , evaluation steps  $k$

Initialize  $V_0$  arbitrarily

**for**  $i = 0, 1, 2, \dots$  **do**

**for**  $j = 1, 2, \dots, k$  **do**

$V_j \leftarrow T^{\pi_i} V_{j-1}$

**end for**

$\pi_{i+1}(s) \leftarrow \arg\max_a [r(s, a) + \gamma \sum_{s'} P(s'|s, a) V_k(s')]$

**end for**

---

## 2.8 Connections to Classical Control Theory

### 2.8.1 Linear Quadratic Regulator (LQR)

For linear dynamics  $s_{t+1} = As_t + Ba_t + w_t$  and quadratic costs  $r(s, a) = -s^T Qs - a^T Ra$ , the optimal value function is quadratic:  $V^*(s) = -s^T Ps$  where  $P$  satisfies the discrete algebraic Riccati equation:

$$P = Q + A^T P A - A^T P B (R + B^T P B)^{-1} B^T P A \quad (2.21)$$

The optimal policy is linear:  $\pi^*(s) = -Ks$  where  $K = (R + B^T P B)^{-1} B^T P A$ .

### 2.8.2 Hamilton-Jacobi-Bellman Equation

For continuous-time systems, the Bellman equation becomes the Hamilton-Jacobi-Bellman (HJB) partial differential equation:

$$\frac{\partial V}{\partial t} + \min_a \left[ r(s, a) + \frac{\partial V}{\partial s} f(s, a) \right] = 0 \quad (2.22)$$

where  $f(s, a)$  is the system dynamics.

## 2.9 Advanced Topics

### 2.9.1 Partially Observable MDPs (POMDPs)

**Definition 2.24** (POMDP). *A POMDP extends an MDP with observations:  $(\mathcal{S}, \mathcal{A}, \mathcal{O}, P, \mathcal{R}, \mathcal{Z}, \gamma)$  where:*

- $\mathcal{O}$  is the observation space
- $\mathcal{Z} : \mathcal{S} \times \mathcal{A} \times \mathcal{O} \rightarrow [0, 1]$  is the observation model

The optimal policy depends on the belief state  $b(s) = \mathbb{P}(S_t = s | h_t)$  where  $h_t$  is the history of observations.

### 2.9.2 Constrained MDPs

**Definition 2.25** (Constrained MDP). *A constrained MDP adds constraint functions  $c_i : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$  and thresholds  $d_i$ :*

$$\max_{\pi} \quad \mathbb{E}^{\pi} \left[ \sum_{t=0}^{\infty} \gamma^t \mathcal{R}(S_t, A_t) \right] \quad (2.23)$$

$$\text{subject to} \quad \mathbb{E}^{\pi} \left[ \sum_{t=0}^{\infty} \gamma^t c_i(S_t, A_t) \right] \leq d_i, \quad i = 1, \dots, m \quad (2.24)$$

Solutions typically use Lagrangian methods and primal-dual algorithms.

## 2.10 Chapter Summary

This chapter established the mathematical foundations of Markov Decision Processes:

- Formal definition of MDPs and regularity assumptions
- Policy and value function theory with existence and uniqueness results
- Bellman equations and optimality conditions
- Contraction mapping theory and convergence guarantees
- Dynamic programming algorithms: value iteration and policy iteration
- Computational complexity analysis
- Connections to classical control theory and advanced extensions

The mathematical framework developed here provides the foundation for all reinforcement learning algorithms. The next chapter examines dynamic programming methods in detail, providing the algorithmic foundation for modern RL techniques.



## Chapter 3

# Dynamic Programming Foundations

Dynamic programming provides the theoretical and algorithmic foundation for reinforcement learning. This chapter develops the mathematical theory of dynamic programming with emphasis on convergence analysis, computational complexity, and connections to classical optimal control.

### 3.1 Principle of Optimality

The fundamental insight underlying dynamic programming is Bellman’s principle of optimality, which enables the decomposition of complex sequential decision problems into simpler subproblems.

**Theorem 3.1** (Principle of Optimality). *An optimal policy has the property that whatever the initial state and initial decision are, the remaining decisions must constitute an optimal policy with regard to the state resulting from the first decision.*

#### 3.1.1 Mathematical Formulation

For an MDP  $(\mathcal{S}, \mathcal{A}, P, \mathcal{R}, \gamma)$ , consider a finite-horizon problem with horizon  $T$ . Define the optimal value function:

$$V_t^*(s) = \max_{\pi} \mathbb{E} \left[ \sum_{k=t}^{T-1} \gamma^{k-t} \mathcal{R}(S_k, A_k) \mid S_t = s, \pi \right] \quad (3.1)$$

**Theorem 3.2** (Finite-Horizon Optimality). *The optimal value function satisfies the*

recursive relation:

$$V_T^*(s) = 0 \quad \forall s \in \mathcal{S} \quad (3.2)$$

$$V_t^*(s) = \max_{a \in \mathcal{A}} \left[ \mathcal{R}(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V_{t+1}^*(s') \right] \quad (3.3)$$

for  $t = T - 1, T - 2, \dots, 0$ .

*Proof.* The proof follows by backward induction. At time  $T$ , no more rewards can be collected, so  $V_T^*(s) = 0$ . For  $t < T$ , any optimal policy must choose the action that maximizes immediate reward plus discounted future value, leading to the recursive formulation.  $\square$

### 3.1.2 Engineering Interpretation

The principle of optimality has direct parallels in engineering optimization:

#### Optimal Control Example

Consider a spacecraft trajectory optimization problem:

- State: position and velocity  $(x, v) \in \mathbb{R}^6$
- Control: thrust vector  $u \in \mathbb{R}^3$
- Dynamics:  $\dot{x} = v, \dot{v} = u/m - \nabla \Phi(x)$  (gravitational field)
- Cost: fuel consumption  $\int_0^T \|u(t)\| dt$

The principle of optimality implies that if we have an optimal trajectory from Earth to Mars, then any sub-trajectory (e.g., from lunar orbit to Mars) must also be optimal for the sub-problem.

## 3.2 Value Iteration: Convergence Analysis

Value iteration is the most fundamental algorithm in dynamic programming, providing a constructive method for computing optimal value functions.

### 3.2.1 Algorithm Description

#### 3.2.2 Convergence Theory

**Theorem 3.3** (Convergence of Value Iteration). *For any initial value function  $V_0$ , the value iteration sequence  $\{V_k\}_{k=0}^\infty$  defined by  $V_{k+1} = T^*V_k$  converges to the unique optimal value function  $V^*$  at geometric rate  $\gamma$ .*

Specifically:

$$\|V_k - V^*\|_\infty \leq \gamma^k \|V_0 - V^*\|_\infty \quad (3.4)$$



---

**Algorithm 3** Value Iteration
 

---

**Require:** MDP  $(\mathcal{S}, \mathcal{A}, P, \mathcal{R}, \gamma)$ , tolerance  $\epsilon > 0$

**Ensure:**  $\epsilon$ -optimal value function  $V$

Initialize  $V_0(s)$  arbitrarily for all  $s \in \mathcal{S}$

$k \leftarrow 0$

**repeat**

**for** each  $s \in \mathcal{S}$  **do**

$V_{k+1}(s) \leftarrow \max_{a \in \mathcal{A}} [\mathcal{R}(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V_k(s')]$

**end for**

$k \leftarrow k + 1$

**until**  $\|V_k - V_{k-1}\|_\infty < \epsilon(1 - \gamma)/(2\gamma)$

**return**  $V_k$

---

*Proof.* Since  $T^*$  is a  $\gamma$ -contraction in the supremum norm and  $V^*$  is the unique fixed point of  $T^*$ , the result follows directly from the Banach fixed point theorem.  $\square$

### 3.2.3 Error Bounds and Stopping Criteria

**Theorem 3.4** (Error Bounds for Value Iteration). *If  $\|V_{k+1} - V_k\|_\infty \leq \delta$ , then:*

$$\|V_k - V^*\|_\infty \leq \frac{\gamma\delta}{1 - \gamma} \quad (3.5)$$

$$\|V_{k+1} - V^*\|_\infty \leq \frac{\delta}{1 - \gamma} \quad (3.6)$$

*Proof.* Using the triangle inequality and contraction property:

$$\|V_k - V^*\|_\infty = \|V_k - T^*V_k + T^*V_k - V^*\|_\infty \quad (3.7)$$

$$\leq \|V_k - T^*V_k\|_\infty + \|T^*V_k - T^*V^*\|_\infty \quad (3.8)$$

$$= \|V_k - V_{k+1}\|_\infty + \gamma\|V_k - V^*\|_\infty \quad (3.9)$$

Solving for  $\|V_k - V^*\|_\infty$  gives the first bound. The second follows similarly.  $\square$

**Corollary 3.5** (Practical Stopping Criterion). *To achieve  $\|V_k - V^*\|_\infty \leq \epsilon$ , it suffices to stop when:*

$$\|V_{k+1} - V_k\|_\infty \leq \epsilon(1 - \gamma) \quad (3.10)$$

### 3.2.4 Computational Complexity

**Theorem 3.6** (Sample Complexity of Value Iteration). *To achieve  $\epsilon$ -optimal value function, value iteration requires:*

$$O\left(\frac{\log(\epsilon^{-1}) + \log(\|V_0 - V^*\|_\infty)}{1 - \gamma}\right) \quad (3.11)$$

*iterations.*

For each iteration:

- **Time complexity:**  $O(|\mathcal{S}|^2|\mathcal{A}|)$  for tabular case
- **Space complexity:**  $O(|\mathcal{S}|)$  for storing value function
- **Total operations:**  $O(|\mathcal{S}|^2|\mathcal{A}| \log(\epsilon^{-1})/(1 - \gamma))$

### 3.3 Policy Iteration: Mathematical Guarantees

Policy iteration alternates between policy evaluation and policy improvement, providing an alternative approach with different computational characteristics.

#### 3.3.1 Policy Evaluation

Given policy  $\pi$ , policy evaluation solves the linear system:

$$V^\pi = T^\pi V^\pi \quad (3.12)$$

In matrix form for finite MDPs:

$$V^\pi = R^\pi + \gamma P^\pi V^\pi \quad (3.13)$$

where  $R^\pi \in \mathbb{R}^{|\mathcal{S}|}$  and  $P^\pi \in \mathbb{R}^{|\mathcal{S}| \times |\mathcal{S}|}$  are policy-specific reward and transition matrices.

**Theorem 3.7** (Unique Solution to Policy Evaluation). *The linear system  $(I - \gamma P^\pi)V^\pi = R^\pi$  has a unique solution:*

$$V^\pi = (I - \gamma P^\pi)^{-1} R^\pi \quad (3.14)$$

since  $\rho(P^\pi) \leq 1$  and  $\gamma < 1$  ensure  $(I - \gamma P^\pi)$  is invertible.

#### 3.3.2 Iterative Policy Evaluation

For large state spaces, direct matrix inversion is computationally prohibitive. Iterative policy evaluation uses:

$$V_{k+1}^\pi = T^\pi V_k^\pi \quad (3.15)$$

**Theorem 3.8** (Convergence of Iterative Policy Evaluation). *The sequence  $\{V_k^\pi\}$  converges geometrically to  $V^\pi$  at rate  $\gamma$ :*

$$\|V_k^\pi - V^\pi\|_\infty \leq \gamma^k \|V_0^\pi - V^\pi\|_\infty \quad (3.16)$$

### 3.3.3 Policy Improvement Analysis

**Theorem 3.9** (Strict Improvement or Optimality). *Given policy  $\pi$  and improved policy  $\pi'$  defined by:*

$$\pi'(s) \in \operatorname{argmax}_{a \in \mathcal{A}} Q^\pi(s, a) \quad (3.17)$$

*Then either:*

1.  $V^{\pi'}(s) > V^\pi(s)$  for some  $s \in \mathcal{S}$  (strict improvement), or
2.  $V^{\pi'}(s) = V^\pi(s)$  for all  $s \in \mathcal{S}$  (optimality)

*Proof.* By construction,  $Q^\pi(s, \pi'(s)) \geq Q^\pi(s, \pi(s)) = V^\pi(s)$  for all  $s$ . If inequality is strict for any state, then by the policy evaluation equations, strict improvement propagates. If equality holds everywhere, then  $\pi$  satisfies the Bellman optimality equation and is optimal.  $\square$

### 3.3.4 Global Convergence

**Theorem 3.10** (Finite Convergence of Policy Iteration). *For finite MDPs, policy iteration converges to an optimal policy in finitely many iterations. Specifically, the number of iterations is bounded by  $|\mathcal{A}|^{|\mathcal{S}|}$ .*

*Proof.* Since each iteration either strictly improves the policy or terminates at optimality, and there are finitely many deterministic policies, convergence must occur in finite time. The bound follows from counting the total number of deterministic policies.  $\square$

## 3.4 Modified Policy Iteration

Modified policy iteration interpolates between value iteration and policy iteration, providing computational flexibility.

### 3.4.1 Algorithm and Convergence

---

#### Algorithm 4 Modified Policy Iteration

---

**Require:** Initial policy  $\pi_0$ , evaluation steps  $m$

$i \leftarrow 0$

**repeat**

$V \leftarrow$  arbitrary initialization

**for**  $k = 1, 2, \dots, m$  **do**

$V \leftarrow T^{\pi_i} V$

**end for**

$\pi_{i+1}(s) \leftarrow \operatorname{argmax}_a [r(s, a) + \gamma \sum_{s'} P(s'|s, a) V(s')]$

$i \leftarrow i + 1$

**until** convergence

---

**Theorem 3.11** (Convergence of Modified Policy Iteration). *Modified policy iteration with  $m \geq 1$  evaluation steps converges to an optimal policy. The convergence rate depends on  $m$ :*

- $m = 1$ : reduces to value iteration with rate  $\gamma$
- $m = \infty$ : reduces to policy iteration with finite convergence
- $1 < m < \infty$ : intermediate convergence rate

### 3.4.2 Optimal Choice of Evaluation Steps

The computational trade-off between evaluation and improvement can be optimized:

**Theorem 3.12** (Optimal Evaluation Steps). *For modified policy iteration, the optimal number of evaluation steps  $m^*$  minimizes total computational cost:*

$$m^* = \operatorname{argmin}_m [\text{cost per iteration} \times \text{number of iterations}] \quad (3.18)$$

*Under reasonable assumptions about computational costs,  $m^* = O(\log(1/(1 - \gamma)))$ .*

## 3.5 Asynchronous Dynamic Programming

Traditional DP algorithms update all states synchronously. Asynchronous variants can offer computational advantages and theoretical insights.

### 3.5.1 Gauss-Seidel Value Iteration

---

**Algorithm 5** Gauss-Seidel Value Iteration

---

```

Order states  $s_1, s_2, \dots, s_n$ 
repeat
  for  $i = 1, 2, \dots, n$  do
     $V(s_i) \leftarrow \max_a \left[ r(s_i, a) + \gamma \sum_j P(s_j | s_i, a) V(s_j) \right]$ 
  end for
until convergence

```

---

**Theorem 3.13** (Convergence of Gauss-Seidel Value Iteration). *Gauss-Seidel value iteration converges to the optimal value function. The convergence rate can be faster than standard value iteration due to more frequent updates.*

### 3.5.2 Prioritized Sweeping

**Definition 3.14** (Bellman Error). *For state  $s$  and value function  $V$ , the Bellman error is:*

$$\delta(s) = \left| \max_a \left[ r(s, a) + \gamma \sum_{s'} P(s' | s, a) V(s') \right] - V(s) \right| \quad (3.19)$$

Prioritized sweeping updates states in order of decreasing Bellman error, focusing computation on states where updates will have the largest impact.

---

**Algorithm 6** Prioritized Sweeping

---

```

Initialize priority queue  $\mathcal{Q}$  with all states
while  $\mathcal{Q}$  not empty do
   $s \leftarrow$  state with highest priority in  $\mathcal{Q}$ 
  Update  $V(s)$  using Bellman equation
  Remove  $s$  from  $\mathcal{Q}$ 
  for each predecessor  $s'$  of  $s$  do
    if Bellman error of  $s'$  exceeds threshold then
      Add  $s'$  to  $\mathcal{Q}$  with updated priority
    end if
  end for
end while

```

---

### 3.5.3 Real-Time Dynamic Programming

Real-time DP focuses updates on states visited by a simulated or actual agent trajectory.

---

**Algorithm 7** Real-Time Dynamic Programming

---

```

Initialize current state  $s$ 
repeat
  Update  $V(s)$  using Bellman equation
  Choose action  $a = \operatorname{argmax}_a Q(s, a)$ 
  Simulate or execute action:  $s \leftarrow s'$  with probability  $P(s'|s, a)$ 
until termination

```

---

**Theorem 3.15** (Convergence of RTDP). *Under appropriate exploration conditions, real-time DP converges to optimal values on the states reachable under the optimal policy.*

## 3.6 Linear Programming Formulation

Dynamic programming problems can be formulated as linear programs, providing alternative solution methods and theoretical insights.

### 3.6.1 Primal LP Formulation

The optimal value function can be found by solving:

$$\text{minimize}_V \quad \sum_{s \in \mathcal{S}} \alpha(s) V(s) \tag{3.20}$$

$$\text{subject to} \quad V(s) \geq r(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V(s') \quad \forall s, a \tag{3.21}$$

where  $\alpha(s) > 0$  represents state weights.

**Theorem 3.16** (LP-DP Equivalence). *The optimal solution to the linear program equals the optimal value function  $V^*$ .*

### 3.6.2 Dual LP Formulation

The dual problem involves finding an optimal state-action visitation measure:

$$\text{maximize}_\mu \quad \sum_{s,a} \mu(s,a)r(s,a) \quad (3.22)$$

$$\text{subject to} \quad \sum_a \mu(s,a) - \gamma \sum_{s',a'} \mu(s',a')P(s|s',a') = \alpha(s) \quad \forall s \quad (3.23)$$

$$\mu(s,a) \geq 0 \quad \forall s,a \quad (3.24)$$

**Theorem 3.17** (Strong Duality). *Under mild conditions, strong duality holds between the primal and dual formulations, and complementary slackness conditions characterize optimal policies.*

## 3.7 Connections to Classical Control Theory

### 3.7.1 Discrete-Time Optimal Control

Consider the discrete-time optimal control problem:

$$\text{minimize} \quad \sum_{t=0}^{T-1} L(x_t, u_t) + L_T(x_T) \quad (3.25)$$

$$\text{subject to} \quad x_{t+1} = f(x_t, u_t) + w_t \quad (3.26)$$

$$u_t \in \mathcal{U}(x_t) \quad (3.27)$$

The dynamic programming solution gives the Hamilton-Jacobi-Bellman equation:

$$V_t(x) = \min_{u \in \mathcal{U}(x)} [L(x, u) + \mathbb{E}[V_{t+1}(f(x, u) + w)]] \quad (3.28)$$

### 3.7.2 Stochastic Optimal Control

For stochastic control systems  $dx_t = f(x_t, u_t)dt + \sigma(x_t, u_t)dW_t$ , the continuous-time HJB equation is:

$$\frac{\partial V}{\partial t} + \min_u \left[ L(x, u) + \frac{\partial V}{\partial x} f(x, u) + \frac{1}{2} \text{tr} \left( \sigma(x, u)^T \frac{\partial^2 V}{\partial x^2} \sigma(x, u) \right) \right] = 0 \quad (3.29)$$

### 3.7.3 Model Predictive Control (MPC)

MPC can be viewed as approximate dynamic programming with receding horizon:

**Algorithm 8** Model Predictive Control**repeat**Measure current state  $x_t$ Solve optimization problem over horizon  $[t, t + H]$ : $u_t^*, \dots, u_{t+H-1}^* = \operatorname{argmin} \sum_{k=0}^{H-1} L(x_{t+k}, u_{t+k}) + L_H(x_{t+H})$ Apply  $u_t^*$  and advance to next time step**until** termination

The connection to DP provides stability and performance guarantees for MPC under appropriate conditions.

## 3.8 Computational Considerations

### 3.8.1 Curse of Dimensionality

The computational complexity of DP algorithms scales exponentially with state space dimension:

- Memory:  $O(|\mathcal{S}|)$  for value function storage
- Computation:  $O(|\mathcal{S}|^2|\mathcal{A}|)$  per iteration
- For continuous spaces: requires discretization or function approximation

### 3.8.2 Approximate Dynamic Programming

To handle large state spaces, approximate DP uses function approximation:

$$V(s) \approx \sum_{i=1}^n w_i \phi_i(s) \quad (3.30)$$

where  $\{\phi_i\}$  are basis functions and  $\{w_i\}$  are parameters.

**Theorem 3.18** (Error Propagation in Approximate DP). *If the approximation error is bounded by  $\epsilon$  in supremum norm:*

$$\|V - \hat{V}\|_\infty \leq \epsilon \quad (3.31)$$

*then the policy derived from  $\hat{V}$  satisfies:*

$$\|V^{\hat{\pi}} - V^*\|_\infty \leq \frac{2\gamma\epsilon}{(1-\gamma)^2} \quad (3.32)$$

## 3.9 Chapter Summary

This chapter developed the mathematical foundations of dynamic programming:

- Principle of optimality and recursive decomposition

- Value iteration: convergence theory, error bounds, complexity analysis
- Policy iteration: linear algebra formulation, finite convergence
- Modified policy iteration and computational trade-offs
- Asynchronous variants: Gauss-Seidel, prioritized sweeping, real-time DP
- Linear programming formulations and duality theory
- Connections to classical optimal control and MPC
- Computational challenges and approximate methods

These algorithmic foundations provide the basis for understanding modern reinforcement learning methods. The next chapter begins our exploration of learning algorithms that estimate value functions from experience rather than exact knowledge of the MDP.



## Chapter 4

# Monte Carlo Methods

Monte Carlo methods form the foundation of model-free reinforcement learning, enabling value function estimation from sample episodes without requiring knowledge of the environment dynamics. This chapter develops the mathematical theory of Monte Carlo estimation in the RL context, with emphasis on convergence analysis and variance reduction techniques.

### 4.1 Monte Carlo Estimation Theory

Monte Carlo methods estimate expectations by sampling. In reinforcement learning, we use sample episodes to estimate value functions without requiring the transition probabilities or reward function.

#### 4.1.1 Basic Monte Carlo Principle

Consider estimating the expectation  $\mathbb{E}[X]$  of random variable  $X$ . The Monte Carlo estimator uses  $n$  independent samples  $X_1, \dots, X_n$ :

$$\hat{\mu}_n = \frac{1}{n} \sum_{i=1}^n X_i \tag{4.1}$$

**Theorem 4.1** (Strong Law of Large Numbers). *If  $\mathbb{E}[|X|] < \infty$ , then  $\hat{\mu}_n \rightarrow \mathbb{E}[X]$  almost surely as  $n \rightarrow \infty$ .*

**Theorem 4.2** (Central Limit Theorem). *If  $\text{Var}(X) = \sigma^2 < \infty$ , then:*

$$\sqrt{n}(\hat{\mu}_n - \mathbb{E}[X]) \xrightarrow{d} \mathcal{N}(0, \sigma^2) \tag{4.2}$$

### 4.1.2 Application to Value Function Estimation

For policy  $\pi$ , the value function is:

$$V^\pi(s) = \mathbb{E}^\pi \left[ \sum_{t=0}^{\infty} \gamma^t R_{t+1} \mid S_0 = s \right] \quad (4.3)$$

Monte Carlo estimation uses sample returns  $G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$  from episodes starting in state  $s$  to estimate  $V^\pi(s)$ .

## 4.2 First-Visit vs. Every-Visit Methods

### 4.2.1 First-Visit Monte Carlo

---

**Algorithm 9** First-Visit Monte Carlo Policy Evaluation

---

**Require:** Policy  $\pi$  to evaluate

Initialize  $V(s) \in \mathbb{R}$  arbitrarily for all  $s \in \mathcal{S}$

Initialize  $Returns(s) \leftarrow$  empty list for all  $s \in \mathcal{S}$

**repeat**

    Generate episode following  $\pi$ :  $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

**for**  $t = T-1, T-2, \dots, 0$  **do**

$G \leftarrow \gamma G + R_{t+1}$

**if**  $S_t$  not appear in  $S_0, S_1, \dots, S_{t-1}$  **then**

            Append  $G$  to  $Returns(S_t)$

$V(S_t) \leftarrow \text{average}(Returns(S_t))$

**end if**

**end for**

**until** convergence

---

### 4.2.2 Every-Visit Monte Carlo

Every-visit MC updates the value estimate every time a state is visited in an episode, not just the first time.

**Theorem 4.3** (Convergence of First-Visit Monte Carlo). *First-visit Monte Carlo converges to  $V^\pi(s)$  as the number of first visits to state  $s$  approaches infinity, assuming:*

1. *Episodes are generated according to policy  $\pi$*
2. *Each state has non-zero probability of being the starting state*
3. *Returns have finite variance*

*Proof.* Each first visit to state  $s$  provides an unbiased sample of the return. By the strong law of large numbers, the sample average converges to the true expectation.  $\square$

**Theorem 4.4** (Convergence of Every-Visit Monte Carlo). *Every-visit Monte Carlo also converges to  $V^\pi(s)$  under similar conditions, despite the correlation between visits within the same episode.*

## 4.3 Variance Reduction Techniques

### 4.3.1 Incremental Implementation

Instead of storing all returns, we can update estimates incrementally:

$$V_{n+1}(s) = V_n(s) + \frac{1}{n+1}[G_n - V_n(s)] \quad (4.4)$$

More generally, with step size  $\alpha$ :

$$V(s) \leftarrow V(s) + \alpha[G - V(s)] \quad (4.5)$$

### 4.3.2 Baseline Subtraction

To reduce variance, we can subtract a baseline  $b(s)$  that doesn't depend on the action:

$$G_t - b(S_t) \quad (4.6)$$

The optimal baseline that minimizes variance is:

$$b^*(s) = \frac{\mathbb{E}[G_t^2 \mid S_t = s]}{\mathbb{E}[G_t \mid S_t = s]} = \mathbb{E}[G_t \mid S_t = s] = V^\pi(s) \quad (4.7)$$

### 4.3.3 Control Variates

For correlated random variable  $Y$  with known expectation  $\mathbb{E}[Y] = \mu_Y$ :

$$\hat{\mu}_{CV} = \hat{\mu}_X - c(\hat{\mu}_Y - \mu_Y) \quad (4.8)$$

The optimal coefficient is:

$$c^* = \frac{\text{Cov}(X, Y)}{\text{Var}(Y)} \quad (4.9)$$

## 4.4 Importance Sampling in RL

Importance sampling enables off-policy learning by weighting samples according to the ratio of target to behavior policy probabilities.

### 4.4.1 Ordinary Importance Sampling

To estimate  $\mathbb{E}_\pi[X]$  using samples from policy  $\mu$ :

$$\hat{\mu}_{IS} = \frac{1}{n} \sum_{i=1}^n \rho_i X_i \quad (4.10)$$

where  $\rho_i = \frac{\pi(A_i|S_i)}{\mu(A_i|S_i)}$  is the importance sampling ratio.

**Theorem 4.5** (Unbiasedness of Importance Sampling).  $\mathbb{E}[\hat{\mu}_{IS}] = \mathbb{E}_\pi[X]$  if  $\mu(a|s) > 0$  whenever  $\pi(a|s) > 0$ .

### 4.4.2 Weighted Importance Sampling

To reduce variance when some importance weights are very large:

$$\hat{\mu}_{WIS} = \frac{\sum_{i=1}^n \rho_i X_i}{\sum_{i=1}^n \rho_i} \quad (4.11)$$

**Theorem 4.6** (Bias-Variance Tradeoff). *Weighted importance sampling is biased but often has lower variance than ordinary importance sampling:*

$$\text{Bias}[\hat{\mu}_{WIS}] \neq 0 \text{ (in general)} \quad (4.12)$$

$$\text{Var}[\hat{\mu}_{WIS}] \leq \text{Var}[\hat{\mu}_{IS}] \text{ (typically)} \quad (4.13)$$

### 4.4.3 Per-Decision Importance Sampling

For episodic tasks, the importance sampling ratio for a complete episode is:

$$\rho_{t:T-1} = \prod_{k=t}^{T-1} \frac{\pi(A_k|S_k)}{\mu(A_k|S_k)} \quad (4.14)$$

This can have very high variance. Per-decision importance sampling uses only the relevant portion of the trajectory.

## 4.5 Off-Policy Monte Carlo Methods

### 4.5.1 Off-Policy Policy Evaluation

### 4.5.2 Off-Policy Monte Carlo Control

---

**Algorithm 10** Off-Policy Monte Carlo Policy Evaluation

---

**Require:** Target policy  $\pi$ , behavior policy  $\mu$ Initialize  $V(s) \in \mathbb{R}$  arbitrarily for all  $s \in \mathcal{S}$ Initialize  $C(s) \leftarrow 0$  for all  $s \in \mathcal{S}$ **repeat**  Generate episode using  $\mu$ :  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$    $G \leftarrow 0$    $W \leftarrow 1$   **for**  $t = T - 1, T - 2, \dots, 0$  **do**     $G \leftarrow \gamma G + R_{t+1}$      $C(S_t) \leftarrow C(S_t) + W$      $V(S_t) \leftarrow V(S_t) + \frac{W}{C(S_t)}[G - V(S_t)]$      $W \leftarrow W \frac{\pi(A_t|S_t)}{\mu(A_t|S_t)}$     **if**  $W = 0$  **then**

break

**end if**  **end for****until** convergence

---

---

**Algorithm 11** Off-Policy Monte Carlo Control

---

Initialize  $Q(s, a) \in \mathbb{R}$  arbitrarily for all  $s, a$ Initialize  $C(s, a) \leftarrow 0$  for all  $s, a$ Initialize  $\pi(s) \leftarrow \operatorname{argmax}_a Q(s, a)$  for all  $s$ **repeat**  Choose any soft policy  $\mu$  (e.g.,  $\epsilon$ -greedy)  Generate episode using  $\mu$    $G \leftarrow 0$    $W \leftarrow 1$   **for**  $t = T - 1, T - 2, \dots, 0$  **do**     $G \leftarrow \gamma G + R_{t+1}$      $C(S_t, A_t) \leftarrow C(S_t, A_t) + W$      $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{W}{C(S_t, A_t)}[G - Q(S_t, A_t)]$      $\pi(S_t) \leftarrow \operatorname{argmax}_a Q(S_t, a)$     **if**  $A_t \neq \pi(S_t)$  **then**

break

**end if**     $W \leftarrow W \frac{1}{\mu(A_t|S_t)}$   **end for****until** convergence

---

## 4.6 Convergence Analysis and Sample Complexity

### 4.6.1 Finite Sample Analysis

**Theorem 4.7** (Finite Sample Bound for Monte Carlo). *Let  $V_n(s)$  be the Monte Carlo estimate after  $n$  visits to state  $s$ . Under bounded rewards  $|R| \leq R_{max}$ :*

$$\mathbb{P}(|V_n(s) - V^\pi(s)| \geq \epsilon) \leq 2 \exp\left(-\frac{2n\epsilon^2(1-\gamma)^2}{R_{max}^2}\right) \quad (4.15)$$

### 4.6.2 Asymptotic Convergence Rate

**Theorem 4.8** (Central Limit Theorem for Monte Carlo). *If  $\text{Var}^\pi[G_t|S_t = s] = \sigma^2(s) < \infty$ , then:*

$$\sqrt{n}(V_n(s) - V^\pi(s)) \xrightarrow{d} \mathcal{N}(0, \sigma^2(s)) \quad (4.16)$$

This gives the convergence rate  $O(n^{-1/2})$ , which is slower than the  $O(n^{-1})$  rate achievable by temporal difference methods under certain conditions.

### 4.6.3 Sample Complexity

**Theorem 4.9** (Sample Complexity of Monte Carlo). *To achieve  $\epsilon$ -accurate value function estimation with probability  $1 - \delta$ :*

$$n \geq \frac{R_{max}^2 \log(2/\delta)}{2\epsilon^2(1-\gamma)^2} \quad (4.17)$$

*samples are sufficient.*

## 4.7 Practical Considerations

### 4.7.1 Exploration vs. Exploitation

Monte Carlo control methods face the exploration-exploitation dilemma. Common approaches:

**Exploring Starts:** Assume episodes start in randomly selected state-action pairs.

**$\epsilon$ -Greedy Policies:** Use soft policies that maintain exploration:

$$\pi(a|s) = \begin{cases} 1 - \epsilon + \frac{\epsilon}{|\mathcal{A}(s)|} & \text{if } a = \operatorname{argmax}_a Q(s, a) \\ \frac{\epsilon}{|\mathcal{A}(s)|} & \text{otherwise} \end{cases} \quad (4.18)$$

### 4.7.2 Function Approximation

For large state spaces, we approximate value functions:

$$V(s) \approx \hat{V}(s, \mathbf{w}) = \mathbf{w}^T \boldsymbol{\phi}(s) \quad (4.19)$$

The Monte Carlo update becomes:

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha[G_t - \hat{V}(S_t, \mathbf{w})]\nabla_{\mathbf{w}}\hat{V}(S_t, \mathbf{w}) \quad (4.20)$$

**Theorem 4.10** (Convergence with Linear Function Approximation). *Under linear function approximation with linearly independent features, Monte Carlo methods converge to the best linear approximation in the  $L^2$  norm weighted by the stationary distribution.*

## 4.8 Chapter Summary

This chapter established the foundations of Monte Carlo methods in reinforcement learning:

- Monte Carlo estimation theory and convergence properties
- First-visit vs. every-visit methods with convergence guarantees
- Variance reduction techniques: baselines, control variates, importance sampling
- Off-policy learning through importance sampling with bias-variance analysis
- Sample complexity bounds and convergence rates
- Practical considerations for exploration and function approximation

Monte Carlo methods provide unbiased estimates and are conceptually simple, but they require complete episodes and have slower convergence than temporal difference methods. The next chapter develops temporal difference learning, which enables learning from individual transitions.





## Chapter 5

# Temporal Difference Learning

Temporal Difference (TD) learning combines ideas from Monte Carlo methods and dynamic programming, enabling learning from incomplete episodes while maintaining the model-free nature of Monte Carlo methods. This chapter develops the mathematical theory of TD learning with emphasis on convergence analysis and the fundamental bias-variance tradeoff.

### 5.1 TD(0) Algorithm and Mathematical Analysis

#### 5.1.1 Basic TD(0) Update

The core insight of temporal difference learning is to use the current estimate of the successor state's value to update the current state's value:

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)] \quad (5.1)$$

The TD error is defined as:

$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t) \quad (5.2)$$

#### 5.1.2 Relationship to Bellman Equation

The TD(0) update can be viewed as a stochastic approximation to the Bellman equation. The expected TD update is:

$$\mathbb{E}[\delta_t | S_t = s] = \mathbb{E}[R_{t+1} + \gamma V(S_{t+1}) - V(S_t) | S_t = s] \quad (5.3)$$

$$= \sum_{s', r} p(s', r | s, \pi(s)) [r + \gamma V(s') - V(s)] \quad (5.4)$$

$$= (T^\pi V)(s) - V(s) \quad (5.5)$$

---

**Algorithm 12** Tabular TD(0) Policy Evaluation
 

---

**Require:** Policy  $\pi$  to evaluate, step size  $\alpha \in (0, 1]$

Initialize  $V(s) \in \mathbb{R}$  arbitrarily for all  $s \in \mathcal{S}$ , except  $V(\text{terminal}) = 0$

**repeat**

    Initialize  $S$

**repeat**

$A \leftarrow$  action given by  $\pi$  for  $S$

        Take action  $A$ , observe  $R, S'$

$V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

**until**  $S$  is terminal

**until** convergence or sufficient accuracy

---

where  $T^\pi$  is the Bellman operator for policy  $\pi$ .

### 5.1.3 Convergence Analysis

**Theorem 5.1** (Convergence of TD(0) - Tabular Case). *For the tabular case with appropriate step size sequence  $\{\alpha_t\}$  satisfying:*

$$\sum_{t=0}^{\infty} \alpha_t = \infty \tag{5.6}$$

$$\sum_{t=0}^{\infty} \alpha_t^2 < \infty \tag{5.7}$$

*TD(0) converges to  $V^\pi$  with probability 1.*

*Proof Sketch.* The proof uses stochastic approximation theory. Define the ODE:

$$\frac{dV}{dt} = \mathbb{E}[\delta_t | V] = T^\pi V - V \tag{5.8}$$

Since  $T^\pi$  is a contraction, the unique fixed point is  $V^\pi$ . The stochastic approximation theorem ensures convergence of the discrete updates to the ODE solution.  $\square$

## 5.2 Bias-Variance Tradeoff in TD Methods

### 5.2.1 Bias Analysis

TD(0) uses the biased estimate  $R_{t+1} + \gamma V(S_{t+1})$  as a target for  $V(S_t)$ , while Monte Carlo uses the unbiased estimate  $G_t$ .

**Theorem 5.2** (Bias of TD Target). *The TD target  $R_{t+1} + \gamma V(S_{t+1})$  has bias:*

$$\text{Bias}[R_{t+1} + \gamma V(S_{t+1})] = \gamma[\hat{V}(S_{t+1}) - V^\pi(S_{t+1})] \tag{5.9}$$

where  $\hat{V}$  is the current estimate.

### 5.2.2 Variance Analysis

**Theorem 5.3** (Variance Comparison). *Under the assumption that value function errors are small, the variance of the TD target is approximately:*

$$\text{Var}[R_{t+1} + \gamma V(S_{t+1})] \approx \text{Var}[R_{t+1}] \quad (5.10)$$

while the Monte Carlo target has variance:

$$\text{Var}[G_t] = \text{Var}\left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1}\right] \quad (5.11)$$

which is typically much larger.

### 5.2.3 Mean Squared Error Decomposition

$$\text{MSE} = \text{Bias}^2 + \text{Variance} + \text{Noise} \quad (5.12)$$

TD methods trade increased bias for reduced variance, often resulting in lower overall MSE and faster convergence.

## 5.3 TD( $\lambda$ ) and Eligibility Traces

TD( $\lambda$ ) provides a family of algorithms that interpolate between TD(0) and Monte Carlo methods through the use of eligibility traces.

### 5.3.1 Forward View: n-step Returns

The n-step return combines rewards from the next n steps with the estimated value of the state reached after n steps:

$$G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n}) \quad (5.13)$$

The n-step TD update is:

$$V(S_t) \leftarrow V(S_t) + \alpha[G_t^{(n)} - V(S_t)] \quad (5.14)$$

### 5.3.2 $\lambda$ -Return

The  $\lambda$ -return combines all n-step returns:

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)} \quad (5.15)$$

**Theorem 5.4** ( $\lambda$ -Return Properties). *The  $\lambda$ -return satisfies:*

$$G_t^\lambda = R_{t+1} + \gamma[(1 - \lambda)V(S_{t+1}) + \lambda G_{t+1}^\lambda] \quad (5.16)$$

$$\lim_{\lambda \rightarrow 0} G_t^\lambda = R_{t+1} + \gamma V(S_{t+1}) \quad (TD(0)) \quad (5.17)$$

$$\lim_{\lambda \rightarrow 1} G_t^\lambda = G_t \quad (\text{Monte Carlo}) \quad (5.18)$$

### 5.3.3 Backward View: Eligibility Traces

Eligibility traces provide an online, incremental implementation of TD( $\lambda$ ):

$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t) \quad (5.19)$$

$$e_t(s) = \begin{cases} \gamma \lambda e_{t-1}(s) + 1 & \text{if } s = S_t \\ \gamma \lambda e_{t-1}(s) & \text{if } s \neq S_t \end{cases} \quad (5.20)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t e_t(s) \quad \forall s \quad (5.21)$$

---

#### Algorithm 13 TD( $\lambda$ ) with Eligibility Traces

---

**Require:** Policy  $\pi$ , step size  $\alpha$ , trace decay  $\lambda$

Initialize  $V(s) \in \mathbb{R}$  arbitrarily for all  $s$

**repeat**

    Initialize  $S$ ,  $e(s) = 0$  for all  $s$

**repeat**

$A \leftarrow$  action given by  $\pi$  for  $S$

        Take action  $A$ , observe  $R, S'$

$\delta \leftarrow R + \gamma V(S') - V(S)$

$e(S) \leftarrow e(S) + 1$

**for all**  $s$  **do**

$V(s) \leftarrow V(s) + \alpha \delta e(s)$

$e(s) \leftarrow \gamma \lambda e(s)$

**end for**

$S \leftarrow S'$

**until**  $S$  is terminal

**until** convergence

---

### 5.3.4 Equivalence Theorem

**Theorem 5.5** (Forward-Backward Equivalence). *Under certain conditions, the forward view (using  $\lambda$ -returns) and backward view (using eligibility traces) produce identical updates when applied offline to a complete episode.*

## 5.4 Convergence Theory for Linear Function Approximation

When the state space is large, we use function approximation:

$$V(s) \approx \hat{V}(s, \mathbf{w}) = \mathbf{w}^T \phi(s) \quad (5.22)$$

The TD(0) update becomes:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [R_{t+1} + \gamma \mathbf{w}_t^T \phi(S_{t+1}) - \mathbf{w}_t^T \phi(S_t)] \phi(S_t) \quad (5.23)$$

### 5.4.1 Projected Bellman Equation

Under linear function approximation, TD(0) converges to the solution of the projected Bellman equation:

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \|\Phi \mathbf{w} - T^\pi(\Phi \mathbf{w})\|_{\mathbf{D}}^2 \quad (5.24)$$

where  $\Phi$  is the feature matrix and  $\mathbf{D}$  is a diagonal matrix of state visitation probabilities.

**Theorem 5.6** (Convergence of Linear TD(0)). *Under linear function approximation, TD(0) converges to:*

$$\mathbf{w}^* = (\Phi^T \mathbf{D} \Phi)^{-1} \Phi^T \mathbf{D} \mathbf{r}^\pi \quad (5.25)$$

where  $\mathbf{r}^\pi$  is the expected reward vector.

### 5.4.2 Error Bounds

**Theorem 5.7** (Approximation Error Bound). *Let  $V^*$  be the optimal value function and  $\hat{V}^*$  be the best linear approximation. Then:*

$$\|V^\pi - \hat{V}^\pi\|_{\mathbf{D}} \leq \frac{1}{1-\gamma} \min_{\mathbf{w}} \|V^\pi - \Phi \mathbf{w}\|_{\mathbf{D}} \quad (5.26)$$

## 5.5 Comparison with Monte Carlo and DP Methods

### 5.5.1 Computational Complexity

Method	Memory	Computation per Step	Episode Completion
DP	$O( \mathcal{S} ^2  \mathcal{A} )$	$O( \mathcal{S} ^2  \mathcal{A} )$	Not Required
MC	$O( \mathcal{S} )$	$O(1)$	Required
TD	$O( \mathcal{S} )$	$O(1)$	Not Required

### 5.5.2 Sample Efficiency

**Theorem 5.8** (Sample Complexity Comparison). *Under certain regularity conditions:*

- TD methods:  $O(\frac{1}{\epsilon^2(1-\gamma)^2})$  samples for  $\epsilon$ -accuracy

- *MC methods*:  $O(\frac{1}{\epsilon^2(1-\gamma)^4})$  samples for  $\epsilon$ -accuracy

TD methods often have better sample efficiency due to lower variance, despite being biased.

### 5.5.3 Bootstrapping vs. Sampling

**Bootstrapping:** Using estimates of successor states (DP, TD) **Sampling:** Using actual experience (MC, TD)

TD methods combine both, leading to:

- Faster learning than MC (bootstrapping)
- Model-free nature (sampling)
- Online learning capability

## 5.6 Advanced Topics

### 5.6.1 Multi-step Methods

The  $n$ -step TD methods generalize between TD(0) and Monte Carlo:

$$V(S_t) \leftarrow V(S_t) + \alpha[G_t^{(n)} - V(S_t)] \quad (5.27)$$

**Theorem 5.9** (Optimal Step Size). *For  $n$ -step methods, there exists an optimal  $n$  that minimizes mean squared error, typically  $n \in [3, 10]$  for many problems.*

### 5.6.2 True Online TD( $\lambda$ )

The classical TD( $\lambda$ ) is not equivalent to the forward view when using function approximation. True online TD( $\lambda$ ) corrects this:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha \delta_t \mathbf{z}_t + \alpha (\mathbf{w}_t^T \phi_t - \mathbf{w}_{t-1}^T \phi_t) (\mathbf{z}_t - \phi_t) \quad (5.28)$$

$$\mathbf{z}_{t+1} = \gamma \lambda \mathbf{z}_t + \phi_{t+1} - \alpha \gamma \lambda (\mathbf{z}_t^T \phi_{t+1}) \phi_{t+1} \quad (5.29)$$

### 5.6.3 Gradient TD Methods

To handle function approximation more rigorously, gradient TD methods minimize the mean squared projected Bellman error:

$$\text{MSPBE}(\mathbf{w}) = \|\Pi(\mathbf{T}^\pi \hat{\mathbf{v}} - \hat{\mathbf{v}})\|_{\mathbf{D}}^2 \quad (5.30)$$

$$\nabla \text{MSPBE}(\mathbf{w}) = 2\Phi^T \mathbf{D}(\Pi(\mathbf{T}^\pi \hat{\mathbf{v}} - \hat{\mathbf{v}})) \quad (5.31)$$

## 5.7 Chapter Summary

This chapter developed the mathematical foundations of temporal difference learning:

- TD(0) algorithm with convergence analysis using stochastic approximation theory
- Bias-variance tradeoff analysis showing TD's advantage in variance reduction
- TD( $\lambda$ ) and eligibility traces providing a spectrum between TD(0) and Monte Carlo
- Convergence theory for linear function approximation with error bounds
- Comparative analysis with Monte Carlo and dynamic programming methods
- Advanced topics including multi-step methods and gradient TD approaches

Temporal difference learning provides the foundation for many modern RL algorithms, combining the best aspects of Monte Carlo and dynamic programming approaches. The next chapter extends these ideas to action-value methods with Q-learning and SARSA.





## Part II

# Function Approximation



---

This part extends reinforcement learning beyond tabular methods to handle large and continuous state spaces through function approximation. We explore advanced temporal difference methods, develop linear function approximation with rigorous convergence analysis, and culminate with neural network approaches that enable deep reinforcement learning.

The transition from exact to approximate methods introduces new challenges including the deadly triad of function approximation, bootstrapping, and off-policy learning. We address these challenges with careful algorithm design and theoretical analysis.

---

## Chapter 6

# Q-Learning and SARSA Extensions

### Chapter Overview

This chapter extends our understanding of temporal difference control by exploring advanced variations of Q-learning and SARSA. We examine multi-step methods, eligibility traces, and theoretical convergence guarantees for off-policy learning. The mathematical analysis includes detailed proofs of convergence conditions and performance bounds.

### From Basic TD to Advanced Control

While Chapter 5 introduced the fundamental concepts of TD learning, real-world applications require more sophisticated approaches. Think of basic Q-learning as learning to drive on a simple track - it works, but for complex scenarios like city driving, you need advanced techniques that can handle delayed rewards, partial observability, and efficient exploration.

## 6.1 Multi-Step Q-Learning

### 6.1.1 n-Step Q-Learning

The basic Q-learning update uses only the immediate next reward and state. Multi-step methods extend this by looking ahead multiple steps:

$$Q_{t+n}(S_t, A_t) = Q_t(S_t, A_t) + \alpha_t [G_{t:t+n} - Q_t(S_t, A_t)] \quad (6.1)$$

where the n-step return is defined as:

$$G_{t:t+n} = R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n \max_a Q_t(S_{t+n}, a) \quad (6.2)$$

---

**Algorithm 14** n-Step Q-Learning
 

---

**Require:** Step size  $\alpha \in (0, 1]$ , small  $\epsilon > 0$ , positive integer  $n$   
 Initialize  $Q(s, a)$  arbitrarily for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ , except  $Q(\text{terminal}, \cdot) = 0$   
 Initialize and store  $S_0$ , select and store an action  $A_0 \sim \pi(\cdot | S_0)$   
**for**  $t = 0, 1, 2, \dots$  **do**  
     **if**  $t < T$  **then**  
         Take action  $A_t$ , observe and store the next reward as  $R_{t+1}$  and the next state as  $S_{t+1}$   
         **if**  $S_{t+1}$  is terminal **then**  
              $T \leftarrow t + 1$   
         **else**  
             Select and store  $A_{t+1} \sim \pi(\cdot | S_{t+1})$   
         **end if**  
     **end if**  
      $\tau \leftarrow t - n + 1$  (the time whose state's estimate is being updated)  
     **if**  $\tau \geq 0$  **then**  
          $G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$   
         **if**  $\tau + n < T$  **then**  
              $G \leftarrow G + \gamma^n \max_a Q(S_{\tau+n}, a)$   
         **end if**  
          $Q(S_\tau, A_\tau) \leftarrow Q(S_\tau, A_\tau) + \alpha[G - Q(S_\tau, A_\tau)]$   
     **end if**  
**end for**

---

### 6.1.2 Theoretical Analysis of n-Step Methods

**Theorem 6.1** (n-Step Q-Learning Convergence). *Under standard conditions (bounded rewards, decreasing step size satisfying  $\sum_t \alpha_t = \infty$  and  $\sum_t \alpha_t^2 < \infty$ , and sufficient exploration), n-step Q-learning converges to the optimal action-value function  $Q^*$  with probability 1.*

*Proof.* The proof follows by showing that the n-step return is an unbiased estimate of the optimal value under the greedy policy, then applying the stochastic approximation convergence theorem.

Let  $\pi_t$  be the greedy policy with respect to  $Q_t$ . The n-step return can be written as:

$$G_{t:t+n} = \mathbb{E}_{\pi_t}[R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n}] \quad (6.3)$$

$$+ \gamma^n \max_a Q_t(S_{t+n}, a) + \text{martingale terms} \quad (6.4)$$

As  $Q_t \rightarrow Q^*$ , the bias in this estimate vanishes, ensuring convergence.  $\square$

## 6.2 Q( $\lambda$ ) Learning

### 6.2.1 Eligibility Traces for Q-Learning

Eligibility traces provide an efficient way to update all state-action pairs based on their recency and frequency of visitation:

$$e_t(s, a) = \begin{cases} \gamma\lambda e_{t-1}(s, a) + 1 & \text{if } s = S_t \text{ and } a = A_t \\ \gamma\lambda e_{t-1}(s, a) & \text{otherwise} \end{cases} \quad (6.5)$$

The Q( $\lambda$ ) update is then:

$$Q_{t+1}(s, a) = Q_t(s, a) + \alpha_t \delta_t e_t(s, a) \quad (6.6)$$

where  $\delta_t = R_{t+1} + \gamma \max_{a'} Q_t(S_{t+1}, a') - Q_t(S_t, A_t)$ .

#### Watkins' Q( $\lambda$ ) vs. Naive Q( $\lambda$ )

There are two main variants of Q( $\lambda$ ):

**Watkins' Q( $\lambda$ ):** Resets eligibility traces when a non-greedy action is taken. **Naive Q( $\lambda$ ):** Does not reset traces, leading to off-policy issues.

Watkins' version maintains the off-policy nature of Q-learning while benefiting from eligibility traces.

## 6.3 SARSA( $\lambda$ ) and True Online Methods

### 6.3.1 SARSA( $\lambda$ ) Algorithm

SARSA( $\lambda$ ) combines the on-policy nature of SARSA with eligibility traces:

### 6.3.2 True Online SARSA( $\lambda$ )

The true online version provides exact equivalence to the forward view:

$$Q_{t+1}(s, a) = Q_t(s, a) + \alpha_t \delta_t^s e_t(s, a) + \alpha_t (Q_t(s, a) - Q_{t-1}(s, a)) (e_t(s, a) - \mathbf{1}_{s,a}(S_t, A_t)) \quad (6.7)$$

where  $\mathbf{1}_{s,a}(S_t, A_t)$  is the indicator function.

---

**Algorithm 15** SARSA( $\lambda$ )
 

---

**Require:** Step size  $\alpha \in (0, 1]$ , trace-decay  $\lambda \in [0, 1]$ , small  $\epsilon > 0$

Initialize  $Q(s, a)$  arbitrarily and  $e(s, a) = 0$  for all  $s, a$

**repeat**

Initialize  $S$ , choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

**repeat**

Take action  $A$ , observe  $R, S'$

Choose  $A'$  from  $S'$  using policy derived from  $Q$

$\delta \leftarrow R + \gamma Q(S', A') - Q(S, A)$

$e(S, A) \leftarrow e(S, A) + 1$

**for** all  $s, a$  **do**

$Q(s, a) \leftarrow Q(s, a) + \alpha \delta e(s, a)$

$e(s, a) \leftarrow \gamma \lambda e(s, a)$

**end for**

$S \leftarrow S'; A \leftarrow A'$

**until**  $S$  is terminal

**until** convergence

---

## 6.4 Double Q-Learning

### 6.4.1 The Maximization Bias Problem

Standard Q-learning suffers from maximization bias due to using the same values for both action selection and evaluation:

#### Understanding Maximization Bias

Imagine you're estimating the value of different investments, but your estimates are noisy. When you always pick the investment with the highest estimated value, you're likely to pick one whose value you've overestimated. This systematic error is maximization bias.

### 6.4.2 Double Q-Learning Algorithm

Double Q-learning maintains two independent value functions  $Q^A$  and  $Q^B$ :

### 6.4.3 Bias Reduction Analysis

**Theorem 6.2** (Double Q-Learning Bias Reduction). *Let  $Q^*(s, a)$  be the true optimal value, and let  $\hat{Q}^A(s, a)$  and  $\hat{Q}^B(s, a)$  be independent unbiased estimators. Then:*

$$\mathbb{E}[\hat{Q}^B(s, \arg \max_a \hat{Q}^A(s, a))] \leq \mathbb{E}[\max_a \hat{Q}^A(s, a)] \quad (6.8)$$

*with equality only when the estimates are deterministic.*



**Algorithm 16** Double Q-Learning

---

**Require:** Step sizes  $\alpha^A, \alpha^B \in (0, 1]$ , small  $\epsilon > 0$   
Initialize  $Q^A(s, a)$  and  $Q^B(s, a)$  arbitrarily for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$   
**repeat**  
    Initialize  $S$   
    **repeat**  
        Choose  $A$  from  $S$  using policy derived from  $Q^A + Q^B$  (e.g.,  $\epsilon$ -greedy)  
        Take action  $A$ , observe  $R, S'$   
        With probability 0.5:  
             $A^* \leftarrow \arg \max_a Q^A(S', a)$   
             $Q^A(S, A) \leftarrow Q^A(S, A) + \alpha^A [R + \gamma Q^B(S', A^*) - Q^A(S, A)]$   
        else:  
             $A^* \leftarrow \arg \max_a Q^B(S', a)$   
             $Q^B(S, A) \leftarrow Q^B(S, A) + \alpha^B [R + \gamma Q^A(S', A^*) - Q^B(S, A)]$   
         $S \leftarrow S'$   
    **until**  $S$  is terminal  
**until** convergence

---

## 6.5 Expected SARSA

### 6.5.1 Algorithm and Convergence

Expected SARSA modifies the SARSA update to use the expected value under the current policy:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma \sum_a \pi(a|S_{t+1}) Q(S_{t+1}, a) - Q(S_t, A_t) \right] \quad (6.9)$$

#### Expected SARSA vs. Q-Learning

Expected SARSA can be viewed as a generalization of both SARSA and Q-learning:

- When  $\pi$  is greedy: Expected SARSA = Q-learning
- When  $\pi$  is the behavior policy: Expected SARSA = SARSA

## 6.6 Performance Analysis and Comparison

### 6.6.1 Sample Complexity Bounds

**Theorem 6.3** (Sample Complexity of Q-Learning with Function Approximation). *For Q-learning with linear function approximation in finite MDPs, the sample complexity to achieve  $\epsilon$ -optimal policy is:*

$$\tilde{O} \left( \frac{d^2 S A}{(1 - \gamma)^4 \epsilon^2} \right) \quad (6.10)$$

where  $d$  is the feature dimension.

## 6.6.2 Empirical Comparison Framework

### Experimental Setup for Algorithm Comparison

Standard benchmarks for comparing TD control algorithms:

1. **Tabular domains:** GridWorld, CliffWalking, Taxi
2. **Function approximation:** Mountain Car, CartPole
3. **Metrics:**
  - Learning curves (reward vs. episodes)
  - Sample efficiency (episodes to threshold)
  - Asymptotic performance
  - Computational cost per update

## 6.7 Advanced Topics

### 6.7.1 Gradient Q-Learning

For continuous action spaces, we can use gradient methods:

$$\theta_{t+1} = \theta_t + \alpha_t \delta_t \nabla_{\theta} Q(S_t, A_t; \theta_t) \quad (6.11)$$

where  $\delta_t = R_{t+1} + \gamma \max_a Q(S_{t+1}, a; \theta_t) - Q(S_t, A_t; \theta_t)$ .

### 6.7.2 Distributional Q-Learning

Instead of learning expected returns, distributional methods learn the full return distribution:

$$Z(s, a) \rightarrow \text{distribution of } G_t \text{ given } S_t = s, A_t = a \quad (6.12)$$

The distributional Bellman equation becomes:

$$Z(s, a) \stackrel{d}{=} R(s, a) + \gamma Z(S', A') \quad (6.13)$$

where  $\stackrel{d}{=}$  denotes equality in distribution.

## 6.8 Implementation Considerations

### 6.8.1 Memory and Computational Efficiency

#### Practical Implementation Tips

1. **Eligibility traces:** Use sparse representations for large state spaces
2. **Experience replay:** Store and reuse past experiences for sample efficiency
3. **Target networks:** Use separate target networks for stable learning
4. **Prioritized updates:** Focus computation on important state-action pairs

### 6.8.2 Hyperparameter Sensitivity

Key hyperparameters and their typical ranges:

- Learning rate  $\alpha$ : Usually 0.01 to 0.5
- Discount factor  $\gamma$ : Typically 0.9 to 0.99
- Trace decay  $\lambda$ : Often 0.9 to 0.95
- Exploration parameter  $\epsilon$ : Start at 1.0, decay to 0.01

## 6.9 Chapter Summary

This chapter extended basic temporal difference learning with advanced techniques that address key limitations:

- **Multi-step methods** balance bias and variance in value estimates
- **Eligibility traces** enable efficient credit assignment over time
- **Double Q-learning** reduces maximization bias in off-policy learning
- **Expected SARSA** provides a unified framework for on-policy and off-policy methods

These extensions are crucial for practical applications and form the foundation for modern deep reinforcement learning algorithms covered in subsequent chapters.

#### Key Takeaways

1. Multi-step methods interpolate between Monte Carlo and one-step TD methods
2. Eligibility traces provide an efficient mechanism for temporal credit assignment
3. Off-policy learning requires careful handling of maximization bias
4. The choice between on-policy and off-policy methods depends on the specific application requirements



## Chapter 7

# Linear Function Approximation

### Chapter Overview

This chapter introduces function approximation to reinforcement learning, focusing on linear methods that enable learning in large or continuous state spaces. We develop the mathematical theory of gradient-based updates, convergence guarantees, and the deadly triad of function approximation. The treatment emphasizes both theoretical understanding and practical implementation considerations.

### Why Function Approximation?

Imagine trying to store a separate value for every possible configuration of a chess board (about  $10^{120}$  states) or every possible sensor reading from a robot ( $\mathbb{R}^n$  continuous space). Tabular methods become impossible. Function approximation allows us to generalize from limited experience to the vast space of possible states by learning a parameterized function.

## 7.1 The Need for Generalization

### 7.1.1 Limitations of Tabular Methods

In tabular reinforcement learning, we maintain explicit tables  $V(s)$  or  $Q(s, a)$  for each state or state-action pair. This approach faces fundamental limitations:

1. **Memory Requirements:**  $O(|\mathcal{S}|)$  for value functions,  $O(|\mathcal{S}| \times |\mathcal{A}|)$  for action-value functions
2. **Learning Speed:** Each state must be visited multiple times
3. **Continuous Spaces:** Infinite state spaces cannot be handled
4. **Generalization:** No sharing of information between similar states

### 7.1.2 Function Approximation Framework

We approximate the value function with a parameterized function:

$$\hat{V}(s; \theta) \approx V^\pi(s) \quad (7.1)$$

or for action-values:

$$\hat{Q}(s, a; \theta) \approx Q^\pi(s, a) \quad (7.2)$$

where  $\theta \in \mathbb{R}^d$  is a parameter vector with  $d \ll |\mathcal{S}|$ .

## 7.2 Linear Function Approximation

### 7.2.1 Feature Representation

In linear function approximation, we represent states using feature vectors:

$$\hat{V}(s; \theta) = \theta^T \phi(s) = \sum_{i=1}^d \theta_i \phi_i(s) \quad (7.3)$$

where  $\phi(s) = [\phi_1(s), \phi_2(s), \dots, \phi_d(s)]^T \in \mathbb{R}^d$  is the feature vector.

#### Feature Engineering Examples

##### GridWorld Features:

- Position coordinates:  $\phi_1(s) = x$ ,  $\phi_2(s) = y$
- Distance to goal:  $\phi_3(s) = \|s - s_{\text{goal}}\|$
- Indicator features:  $\phi_i(s) = \mathbf{1}[s = s_i]$  (one-hot encoding)

##### CartPole Features:

- State variables:  $\phi_1(s) = \text{position}$ ,  $\phi_2(s) = \text{velocity}$
- Polynomial features:  $\phi_3(s) = \text{position}^2$ ,  $\phi_4(s) = \text{position} \times \text{velocity}$

### 7.2.2 Linear Action-Value Approximation

For control problems, we approximate action-value functions:

$$\hat{Q}(s, a; \theta) = \theta^T \phi(s, a) \quad (7.4)$$

where  $\phi(s, a)$  can be constructed as:

- **Separate features:**  $\phi(s, a) = [\phi^{(1)}(s, a), \phi^{(2)}(s, a), \dots]^T$
- **State-action concatenation:**  $\phi(s, a) = [\phi_s(s), \phi_a(a)]^T$
- **Tile coding:** Discretize continuous spaces with overlapping tiles

## 7.3 Gradient-Based Learning

### 7.3.1 Value Function Approximation with Gradient Descent

We minimize the mean squared error between our approximation and target values:

$$J(\theta) = \mathbb{E}_{\mu} \left[ \left( V^{\pi}(s) - \hat{V}(s; \theta) \right)^2 \right] \quad (7.5)$$

where  $\mu$  is a state distribution.

The gradient descent update is:

$$\theta_{t+1} = \theta_t - \frac{1}{2} \alpha \nabla_{\theta} J(\theta_t) \quad (7.6)$$

For linear approximation:

$$\nabla_{\theta} \hat{V}(s; \theta) = \phi(s) \quad (7.7)$$

### 7.3.2 Stochastic Gradient Descent

Since we don't know  $V^{\pi}(s)$ , we use sample-based updates. For a sample  $(S_t, V_t)$  where  $V_t$  is our target:

$$\theta_{t+1} = \theta_t + \alpha \left[ V_t - \hat{V}(S_t; \theta_t) \right] \phi(S_t) \quad (7.8)$$

---

**Algorithm 17** Gradient Monte Carlo for Value Approximation

---

**Require:** Differentiable function  $\hat{V}(s; \theta)$ , policy  $\pi$

Initialize  $\theta$  arbitrarily

**repeat**

Generate an episode  $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$  following  $\pi$

**for**  $t = 0, 1, \dots, T - 1$  **do**

$G_t \leftarrow$  return from time  $t$

$\theta \leftarrow \theta + \alpha [G_t - \hat{V}(S_t; \theta)] \nabla_{\theta} \hat{V}(S_t; \theta)$

**end for**

**until** convergence

---

## 7.4 Temporal Difference with Function Approximation

### 7.4.1 Semi-Gradient TD(0)

For temporal difference learning, we use the current estimate of the successor state as the target:

$$\theta_{t+1} = \theta_t + \alpha \left[ R_{t+1} + \gamma \hat{V}(S_{t+1}; \theta_t) - \hat{V}(S_t; \theta_t) \right] \phi(S_t) \quad (7.9)$$

This is called "semi-gradient" because we don't take the gradient with respect to  $\theta$  in the target  $R_{t+1} + \gamma \hat{V}(S_{t+1}; \theta_t)$ .

---

**Algorithm 18** Semi-gradient TD(0) for Value Approximation

---

**Require:** Differentiable function  $\hat{V}(s; \theta)$ , policy  $\pi$   
Initialize  $\theta$  arbitrarily  
**repeat**  
    Initialize  $S$   
    **repeat**  
         $A \leftarrow$  action given by  $\pi$  for  $S$   
        Take action  $A$ , observe  $R, S'$   
         $\theta \leftarrow \theta + \alpha[R + \gamma\hat{V}(S'; \theta) - \hat{V}(S; \theta)]\nabla_{\theta}\hat{V}(S; \theta)$   
         $S \leftarrow S'$   
    **until**  $S$  is terminal  
**until** convergence

---

### 7.4.2 Q-Learning with Function Approximation

For action-value approximation:

$$\theta_{t+1} = \theta_t + \alpha \left[ R_{t+1} + \gamma \max_a \hat{Q}(S_{t+1}, a; \theta_t) - \hat{Q}(S_t, A_t; \theta_t) \right] \nabla_{\theta} \hat{Q}(S_t, A_t; \theta_t) \quad (7.10)$$

## 7.5 Convergence Analysis

### 7.5.1 The Projection Matrix

For linear function approximation, define the projection matrix:

$$\mathbf{P} = \Phi(\Phi^T \mathbf{D} \Phi)^{-1} \Phi^T \mathbf{D} \quad (7.11)$$

where  $\Phi$  is the  $|\mathcal{S}| \times d$  feature matrix and  $\mathbf{D}$  is a diagonal matrix of state visitation probabilities.

### 7.5.2 Fixed Point Analysis

**Theorem 7.1** (Convergence of Linear TD(0)). *For linear function approximation with features  $\phi(s)$ , semi-gradient TD(0) converges to the fixed point:*

$$\theta_{TD} = (\Phi^T \mathbf{D} \Phi)^{-1} \Phi^T \mathbf{D} V^{\pi} \quad (7.12)$$

*This represents the best linear approximation to  $V^{\pi}$  in the weighted  $L_2$  norm with weights given by the state distribution.*

*Proof.* The TD update can be written as:

$$\theta_{t+1} = \theta_t + \alpha(\mathbf{R} + \gamma \mathbf{P} \Phi \theta_t - \Phi \theta_t)^T \mathbf{D} \Phi \quad (7.13)$$



At the fixed point:  $\Phi\theta_{TD} = \mathbf{P}(\mathbf{R} + \gamma\mathbf{P}\Phi\theta_{TD})$

Solving for  $\theta_{TD}$  yields the stated result.  $\square$

### 7.5.3 Mean Squared Error Bound

**Theorem 7.2** (TD Error Bound). *The steady-state error of linear TD(0) satisfies:*

$$\|\hat{V} - V^\pi\|_\mu^2 \leq \frac{1}{1 - \gamma^2} \min_{\theta} \|V^\pi - \Phi\theta\|_\mu^2 \quad (7.14)$$

where  $\|\cdot\|_\mu$  is the weighted  $L_2$  norm under distribution  $\mu$ .

## 7.6 The Deadly Triad

### 7.6.1 Instability in Function Approximation

The combination of three factors can lead to instability:

1. **Function approximation:** Using parameterized functions instead of tables
2. **Bootstrapping:** Using current estimates to update current estimates (as in TD learning)
3. **Off-policy learning:** Learning about a different policy than the one generating data

#### Baird's Counterexample

Baird (1995) constructed a simple MDP where off-policy TD learning with linear function approximation diverges:

States:  $\{s_1, s_2, \dots, s_7\}$  with terminal state  $s_7$  Features:  $\phi(s_i) = [2\mathbf{1}[i \leq 6], \mathbf{1}[i = 1], \dots, \mathbf{1}[i = 6]]^T$  for  $i = 1, \dots, 6$  Behavior policy: Uniform random among available actions Target policy: Always take "dashed" action leading to  $s_7$

Despite having a unique solution to the projected Bellman equation, the semi-gradient algorithm diverges.

### 7.6.2 Gradient-TD Methods

To address instability, gradient-TD methods perform true gradient descent on the mean squared projected Bellman error:

$$\text{MSPBE}(\theta) = \|\Pi T^\pi \hat{V}_\theta - \hat{V}_\theta\|_\mu^2 \quad (7.15)$$

The gradient-TD update is:

$$\theta_{t+1} = \theta_t + \alpha(\delta_t \phi_t - \gamma \phi_{t+1} \phi_t^T w_t) \quad (7.16)$$

where  $w_t$  is an auxiliary parameter vector.

## 7.7 Least-Squares Methods

### 7.7.1 Least-Squares Temporal Difference (LSTD)

LSTD computes the TD fixed point directly by solving:

$$\mathbf{A}\theta = \mathbf{b} \quad (7.17)$$

where:

$$\mathbf{A} = \sum_{t=1}^T \phi(S_t)(\phi(S_t) - \gamma\phi(S_{t+1}))^T \quad (7.18)$$

$$\mathbf{b} = \sum_{t=1}^T \phi(S_t)R_{t+1} \quad (7.19)$$

---

#### Algorithm 19 LSTD for Value Approximation

---

**Require:** Features  $\phi(s)$ , policy  $\pi$ , regularization  $\epsilon > 0$

Initialize  $\mathbf{A} = \epsilon \mathbf{I}$ ,  $\mathbf{b} = \mathbf{0}$

**repeat**

Collect episode following  $\pi$

**for** each transition  $(S_t, R_{t+1}, S_{t+1})$  **do**

$\mathbf{A} \leftarrow \mathbf{A} + \phi(S_t)(\phi(S_t) - \gamma\phi(S_{t+1}))^T$

$\mathbf{b} \leftarrow \mathbf{b} + \phi(S_t)R_{t+1}$

**end for**

$\theta \leftarrow \mathbf{A}^{-1}\mathbf{b}$

**until** convergence

---

### 7.7.2 Recursive Least-Squares Implementation

To avoid matrix inversion, use the Sherman-Morrison formula:

$$\mathbf{A}_{t+1}^{-1} = \mathbf{A}_t^{-1} - \frac{\mathbf{A}_t^{-1}\phi_t\phi_t^T\mathbf{A}_t^{-1}}{1 + \phi_t^T\mathbf{A}_t^{-1}\phi_t} \quad (7.20)$$

## 7.8 Feature Construction

### 7.8.1 Tile Coding

Tile coding provides a way to discretize continuous spaces with overlapping receptive fields:

- Divide the state space into multiple overlapping grids (tilings)
- For state  $s$ , activate one tile from each tiling

- Feature vector has one 1 for each active tile, 0 elsewhere

#### Tile Coding for Mountain Car

For Mountain Car with position  $p \in [-1.2, 0.6]$  and velocity  $v \in [-0.07, 0.07]$ :

- Create 8 tilings, each  $8 \times 8$  grid
- Offset each tiling by  $(i \cdot 0.225/8, i \cdot 0.01575/8)$  for tiling  $i$
- Total features:  $8 \times 8 \times 8 = 512$
- Active features per state: 8 (one per tiling)

### 7.8.2 Radial Basis Functions

RBFs use Gaussian-like functions centered at prototypes:

$$\phi_i(s) = \exp\left(-\frac{\|s - c_i\|^2}{2\sigma_i^2}\right) \quad (7.21)$$

where  $c_i$  are prototype locations and  $\sigma_i$  are width parameters.

### 7.8.3 Fourier Basis

For states  $s \in [0, 1]^n$ , Fourier features are:

$$\phi_i(s) = \cos(\pi c_i^T s) \quad (7.22)$$

where  $c_i$  are integer vectors determining the frequency and phase.

## 7.9 Policy Gradient with Linear Approximation

### 7.9.1 Linear Policy Approximation

For discrete actions, use softmax policy:

$$\pi(a|s; \theta) = \frac{e^{\theta^T \phi(s,a)}}{\sum_{a'} e^{\theta^T \phi(s,a')}} \quad (7.23)$$

For continuous actions, use Gaussian policy:

$$\pi(a|s; \theta) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(a - \theta^T \phi(s))^2}{2\sigma^2}\right) \quad (7.24)$$

### 7.9.2 REINFORCE with Linear Features

The policy gradient for linear policies has a simple form:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi} [G_t \nabla_{\theta} \log \pi(A_t | S_t; \theta)] \quad (7.25)$$

For the softmax policy:

$$\nabla_{\theta} \log \pi(a | s; \theta) = \phi(s, a) - \sum_{a'} \pi(a' | s; \theta) \phi(s, a') \quad (7.26)$$

## 7.10 Practical Considerations

### 7.10.1 Feature Scaling and Normalization

#### Feature Engineering Guidelines

1. **Normalization:** Scale features to similar ranges (e.g.,  $[0, 1]$  or  $[-1, 1]$ )
2. **Standardization:** Zero mean, unit variance for each feature
3. **Avoid redundancy:** Remove linearly dependent features
4. **Incremental features:** Add polynomial or interaction terms carefully

### 7.10.2 Regularization Techniques

To prevent overfitting:

- **L2 regularization:** Add  $\lambda \|\theta\|^2$  to the objective
- **L1 regularization:** Add  $\lambda \|\theta\|_1$  for sparsity
- **Early stopping:** Monitor validation performance
- **Feature selection:** Use domain knowledge or automated methods

## 7.11 Chapter Summary

Linear function approximation provides the foundation for learning in large state spaces:

- **Representation:** Linear combination of hand-crafted features
- **Learning:** Gradient-based updates with convergence guarantees
- **Challenges:** The deadly triad can cause instability
- **Solutions:** Gradient-TD methods and least-squares approaches

The principles developed here extend naturally to neural networks, which we examine in the next chapter.

**Key Takeaways**

1. Function approximation enables generalization across similar states
2. Linear methods provide strong theoretical guarantees and interpretability
3. Feature engineering is crucial for performance in linear methods
4. The deadly triad poses fundamental challenges that require careful algorithm design
5. Gradient-based methods form the foundation for modern deep RL



## Chapter 8

# Neural Networks in Reinforcement Learning

### Chapter Overview

This chapter introduces deep reinforcement learning by extending function approximation to neural networks. We cover the fundamental algorithms (DQN, Double DQN, Dueling DQN), analyze the unique challenges of combining deep learning with RL, and provide mathematical foundations for understanding stability and convergence in neural network-based RL systems.

### From Linear to Nonlinear

While linear function approximation works well when good features are available, many real-world problems require learning complex, nonlinear relationships. Neural networks provide this capability automatically - they learn both the features and the mapping from features to values. Think of neural networks as universal function approximators that can discover the right representation for any given problem.

## 8.1 Neural Network Fundamentals for RL

### 8.1.1 Universal Function Approximation

Neural networks can approximate any continuous function to arbitrary accuracy given sufficient capacity:

**Theorem 8.1** (Universal Approximation Theorem). *Let  $\sigma$  be a non-constant, bounded, and continuous activation function. Then finite sums of the form:*

$$f(x) = \sum_{i=1}^N v_i \sigma(w_i^T x + b_i) \quad (8.1)$$

are dense in  $C(K)$ , the space of continuous functions on any compact subset  $K \subset \mathbb{R}^n$ .

For RL, this means neural networks can theoretically represent any value function or policy.

### 8.1.2 Deep Neural Network Architecture

A deep neural network with  $L$  layers computes:

$$h^{(0)} = x \tag{8.2}$$

$$h^{(l)} = \sigma(W^{(l)}h^{(l-1)} + b^{(l)}) \quad \text{for } l = 1, \dots, L-1 \tag{8.3}$$

$$f(x; \theta) = W^{(L)}h^{(L-1)} + b^{(L)} \tag{8.4}$$

where  $\theta = \{W^{(l)}, b^{(l)}\}_{l=1}^L$  represents all parameters.

### 8.1.3 Gradient Computation via Backpropagation

The gradient with respect to parameters is computed using the chain rule:

$$\frac{\partial f}{\partial W^{(l)}} = \frac{\partial f}{\partial h^{(l)}} \frac{\partial h^{(l)}}{\partial W^{(l)}} \tag{8.5}$$

For RL, we typically minimize:

$$L(\theta) = \mathbb{E} \left[ (Y - f(S; \theta))^2 \right] \tag{8.6}$$

where  $Y$  is the target (e.g., TD target or Monte Carlo return).

## 8.2 Deep Q-Networks (DQN)

### 8.2.1 The DQN Algorithm

DQN extends Q-learning to neural networks with two key innovations:

1. **Experience Replay:** Store transitions in a replay buffer and sample mini-batches
2. **Target Network:** Use a separate, slowly updated network for computing targets

### 8.2.2 Experience Replay Analysis

Experience replay addresses several issues:

- **Sample Efficiency:** Each experience can be used multiple times
- **Correlation:** Breaking temporal correlations in training data
- **Distribution:** Smoothing over many past behaviors



**Algorithm 20** Deep Q-Network (DQN)**Require:** Replay buffer capacity  $N$ , mini-batch size  $m$ , target update frequency  $C$ 


---

```

Initialize replay buffer  $\mathcal{D}$  with capacity  $N$ 
Initialize action-value function  $Q$  with random weights  $\theta$ 
Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$ 
for episode = 1,  $M$  do
  Initialize sequence  $s_1$ 
  for  $t = 1, T$  do
    With probability  $\epsilon$  select random action  $a_t$ 
    otherwise select  $a_t = \arg \max_a Q(s_t, a; \theta)$ 
    Execute action  $a_t$  and observe reward  $r_t$  and state  $s_{t+1}$ 
    Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $\mathcal{D}$ 
    Sample random mini-batch of transitions from  $\mathcal{D}$ 
    Set  $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(s_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$ 
    Perform gradient descent step on  $(y_j - Q(s_j, a_j; \theta))^2$ 
    Every  $C$  steps reset  $\hat{Q} = Q$  (i.e.,  $\theta^- = \theta$ )
  end for
end for

```

---

**Theorem 8.2** (Experience Replay Convergence). *Under standard assumptions and with appropriate replay buffer management, DQN with experience replay converges to a neighborhood of the optimal  $Q$ -function. The convergence rate depends on the replay buffer size and sampling strategy.*

### 8.2.3 Target Network Stabilization

The target network addresses the non-stationarity of the learning target:

Without target network:  $y_t = r_t + \gamma \max_a Q(s_{t+1}, a; \theta_t)$  With target network:  $y_t = r_t + \gamma \max_a Q(s_{t+1}, a; \theta^-)$

where  $\theta^-$  is updated less frequently than  $\theta$ .

#### Target Network Update Strategies

**Hard Update:**  $\theta^- \leftarrow \theta$  every  $C$  steps **Soft Update:**  $\theta^- \leftarrow \tau\theta + (1 - \tau)\theta^-$  with  $\tau \ll 1$   
**Polyak Averaging:** Exponential moving average of parameters

## 8.3 Advanced DQN Variants

### 8.3.1 Double DQN

Double DQN addresses maximization bias by decoupling action selection and evaluation:

$$y_t = r_t + \gamma Q(s_{t+1}, \arg \max_a Q(s_{t+1}, a; \theta_t); \theta_t^-) \quad (8.7)$$

---

**Algorithm 21** Double DQN Update
 

---

```

Sample mini-batch from replay buffer
for each transition  $(s, a, r, s')$  in mini-batch do
     $a^* \leftarrow \arg \max_{a'} Q(s', a'; \theta)$  {Online network selects action}
     $y \leftarrow r + \gamma Q(s', a^*; \theta^-)$  {Target network evaluates action}
    Compute loss:  $L = (y - Q(s, a; \theta))^2$ 
end for
Update  $\theta$  using gradient of total loss
    
```

---

### 8.3.2 Dueling DQN

Dueling DQN separates state value and advantage estimation:

$$Q(s, a; \theta) = V(s; \theta_v) + A(s, a; \theta_a) - \frac{1}{|\mathcal{A}|} \sum_{a'} A(s, a'; \theta_a) \quad (8.8)$$

where:

- $V(s; \theta_v)$  is the state value function
- $A(s, a; \theta_a)$  is the advantage function
- The subtraction ensures identifiability

#### Why Dueling Architecture?

In many states, the choice of action doesn't significantly affect the value. For example, in Atari games, during periods where no enemies are visible, all actions might have similar values. The dueling architecture allows the network to learn when action selection matters and when it doesn't.

### 8.3.3 Prioritized Experience Replay

Instead of uniform sampling, prioritize experiences based on their TD error:

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha} \quad (8.9)$$

where  $p_i = |\delta_i| + \epsilon$  and  $\alpha$  controls the prioritization strength.

To correct for the bias introduced by non-uniform sampling, use importance sampling weights:

$$w_i = \left( \frac{1}{N \cdot P(i)} \right)^\beta \quad (8.10)$$

## 8.4 Policy-Based Deep RL

### 8.4.1 Deep Policy Networks

For continuous action spaces, parameterize the policy directly:

**Deterministic Policy:**

$$a = \mu(s; \theta) \quad (8.11)$$

**Stochastic Policy:**

$$\pi(a|s; \theta) = \mathcal{N}(a; \mu(s; \theta_\mu), \sigma^2(s; \theta_\sigma)) \quad (8.12)$$

### 8.4.2 Deep Deterministic Policy Gradient (DDPG)

DDPG extends DQN to continuous actions using an actor-critic approach:

**Actor Update:**

$$\nabla_{\theta_\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s_i, a; \theta_Q) \big|_{a=\mu(s_i)} \nabla_{\theta_\mu} \mu(s_i; \theta_\mu) \quad (8.13)$$

**Critic Update:**

$$y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}; \theta'_\mu); \theta'_Q) \quad (8.14)$$

---

**Algorithm 22** Deep Deterministic Policy Gradient (DDPG)

---

Initialize critic  $Q(s, a; \theta_Q)$  and actor  $\mu(s; \theta_\mu)$  with random weights  
Initialize target networks  $Q'$  and  $\mu'$  with weights  $\theta'_Q \leftarrow \theta_Q$ ,  $\theta'_\mu \leftarrow \theta_\mu$   
Initialize replay buffer  $\mathcal{R}$   
**for** episode = 1,  $M$  **do**  
    Initialize random process  $\mathcal{N}$  for action exploration  
    Receive initial observation state  $s_1$   
    **for**  $t = 1, T$  **do**  
        Select action  $a_t = \mu(s_t; \theta_\mu) + \mathcal{N}_t$   
        Execute action  $a_t$  and observe reward  $r_t$  and new state  $s_{t+1}$   
        Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $\mathcal{R}$   
        Sample random mini-batch of  $N$  transitions from  $\mathcal{R}$   
        Set  $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}; \theta'_\mu); \theta'_Q)$   
        Update critic by minimizing loss:  $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i; \theta_Q))^2$   
        Update actor using sampled policy gradient  
        Update target networks:  $\theta'_Q \leftarrow \tau \theta_Q + (1 - \tau) \theta'_Q$ ,  $\theta'_\mu \leftarrow \tau \theta_\mu + (1 - \tau) \theta'_\mu$   
    **end for**  
**end for**

---

## 8.5 Challenges in Deep RL

### 8.5.1 Sample Efficiency

Deep RL typically requires millions of samples due to:

- High-dimensional parameter spaces
- Sparse and delayed rewards
- Exploration challenges
- Non-stationarity of the learning problem

### 8.5.2 Stability and Convergence

Neural networks in RL face unique stability challenges:

1. **Moving Targets:** Both the policy and value function change during learning
2. **Correlated Data:** Sequential observations violate i.i.d. assumptions
3. **Deadly Triad:** Function approximation + bootstrapping + off-policy learning

#### The Overestimation Problem

Q-learning with function approximation often overestimates action values due to:

- Maximization bias in the Bellman update
- Function approximation errors
- Limited exploration leading to incomplete information

Double DQN partially addresses this, but the problem persists in complex environments.

### 8.5.3 Hyperparameter Sensitivity

Deep RL algorithms are notoriously sensitive to hyperparameters:

- Learning rates for actor and critic
- Network architecture choices
- Replay buffer size and sampling strategy
- Target network update frequency
- Exploration schedule

## 8.6 Advanced Architectures

### 8.6.1 Recurrent Neural Networks for Partial Observability

For partially observable environments, use LSTM or GRU networks:

$$h_t = \text{LSTM}(h_{t-1}, s_t) \quad (8.15)$$

$$Q(s_t, a; \theta) = f(h_t; \theta) \quad (8.16)$$

The hidden state  $h_t$  maintains information about the history of observations.

### 8.6.2 Attention Mechanisms

For environments with complex spatial or temporal structure:

$$\alpha_{t,i} = \frac{\exp(e_{t,i})}{\sum_{j=1}^T \exp(e_{t,j})} \quad (8.17)$$

$$c_t = \sum_{i=1}^T \alpha_{t,i} h_i \quad (8.18)$$

where  $e_{t,i}$  is the attention energy between current time  $t$  and history time  $i$ .

### 8.6.3 Graph Neural Networks

For problems with graph structure (e.g., molecular design, social networks):

$$h_v^{(l+1)} = \sigma \left( W^{(l)} h_v^{(l)} + \sum_{u \in \mathcal{N}(v)} h_u^{(l)} \right) \quad (8.19)$$

where  $\mathcal{N}(v)$  are the neighbors of node  $v$ .

## 8.7 Regularization and Generalization

### 8.7.1 Dropout in RL Networks

Apply dropout during training to prevent overfitting:

$$h_{\text{dropout}}^{(l)} = h^{(l)} \odot m \quad (8.20)$$

where  $m$  is a binary mask with probability  $p$  of being 1.

### 8.7.2 Batch Normalization

Normalize activations to stabilize training:

$$\hat{h}^{(l)} = \frac{h^{(l)} - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \quad (8.21)$$

However, batch normalization can interfere with RL dynamics and requires careful application.

### 8.7.3 Weight Decay and L2 Regularization

Add regularization term to the loss:

$$L_{\text{total}} = L_{\text{RL}} + \lambda \sum_l \|W^{(l)}\|_2^2 \quad (8.22)$$

## 8.8 Theoretical Analysis

### 8.8.1 Neural Tangent Kernel Theory

For infinitely wide networks, the dynamics during training can be characterized by the Neural Tangent Kernel:

$$\Theta(x, x') = \langle \nabla_{\theta} f(x; \theta_0), \nabla_{\theta} f(x'; \theta_0) \rangle \quad (8.23)$$

This provides theoretical insights into trainability and generalization.

### 8.8.2 Approximation Error Analysis

**Theorem 8.3** (Deep Network Approximation). *For functions with compositional structure, deep networks can achieve exponentially better approximation rates than shallow networks. Specifically, for functions with  $d$ -dimensional input and smoothness  $s$ , the approximation error scales as:*

$$\epsilon \sim n^{-s/d} \quad (8.24)$$

where  $n$  is the number of parameters.

## 8.9 Implementation Considerations

### 8.9.1 Computational Graph and Automatic Differentiation

Modern deep learning frameworks construct computational graphs:

- **Forward pass:** Compute function values
- **Backward pass:** Compute gradients via reverse-mode AD
- **Parameter updates:** Apply optimization algorithms

### 8.9.2 Hardware Acceleration

Deep RL benefits significantly from:

- **GPU acceleration:** Parallel matrix operations
- **Mixed precision:** Using float16 for forward pass, float32 for gradients
- **Model parallelism:** Distributing large networks across devices
- **Data parallelism:** Processing multiple environments simultaneously

### 8.9.3 Memory Management

Efficient memory usage is crucial:

- **Gradient checkpointing:** Trade computation for memory
- **Experience replay optimization:** Efficient storage and sampling
- **Model compression:** Pruning and quantization techniques

## 8.10 Debugging and Diagnostics

### 8.10.1 Common Deep RL Failures

#### Debugging Checklist

##### Learning Curves:

- Flat learning: Check exploration, learning rates, reward scaling
- Instability: Examine target network updates, replay buffer size
- Slow convergence: Verify network capacity, feature normalization

##### Network Diagnostics:

- Gradient norms: Should be neither too small nor too large
- Activation distributions: Avoid saturation in nonlinear units
- Weight magnitudes: Monitor for explosion or vanishing

### 8.10.2 Ablation Studies

Systematically remove components to understand their importance:

- Target networks vs. single network
- Experience replay vs. online learning
- Different network architectures
- Various regularization techniques

## 8.11 Chapter Summary

Neural networks extend the reach of reinforcement learning to complex, high-dimensional problems:

- **Universal approximation:** Neural networks can represent complex value functions and policies
- **Key algorithms:** DQN, Double DQN, Dueling DQN provide foundation for deep RL
- **Stability challenges:** The deadly triad requires careful algorithm design
- **Practical considerations:** Hardware, memory, and debugging are crucial for success

The combination of neural networks with RL has enabled breakthroughs in game playing, robotics, and control, but requires careful attention to the unique challenges of this combination.

#### Key Takeaways

1. Neural networks provide the representational power needed for complex RL problems
2. Experience replay and target networks are crucial for stability
3. Deep RL combines challenges from both deep learning and reinforcement learning
4. Careful algorithm design and implementation are essential for success
5. Understanding both the theoretical foundations and practical considerations is necessary



## Part III

# Policy Methods



---

This part develops policy-based reinforcement learning methods that directly optimize policies rather than value functions. We begin with policy gradient methods and the policy gradient theorem, develop actor-critic algorithms that combine policy and value function learning, and conclude with advanced policy optimization techniques.

Policy methods are particularly important for continuous control problems and provide the foundation for many state-of-the-art algorithms. We emphasize both theoretical understanding and practical implementation considerations.

---

## Chapter 9

# Policy Gradient Methods - Summary

### Chapter Overview

Policy gradient methods directly optimize parameterized policies using gradient ascent on expected return. This chapter would cover the policy gradient theorem, REINFORCE algorithm, variance reduction techniques, and natural policy gradients.

Key topics include:

- Policy gradient theorem and mathematical foundations
- REINFORCE algorithm and baseline methods
- Natural policy gradients and trust regions
- Variance reduction techniques
- Continuous action space handling



## Chapter 10

# Actor-Critic Methods - Summary

### Chapter Overview

Actor-critic methods combine the benefits of policy gradient methods (actor) with value function estimation (critic). This chapter would cover basic actor-critic, A2C, A3C, and advanced variants.

Key topics include:

- Actor-critic architecture and mathematical framework
- Advantage Actor-Critic (A2C) and Asynchronous A3C
- Generalized Advantage Estimation (GAE)
- Off-policy actor-critic methods
- Practical implementation considerations





## Chapter 11

# Advanced Policy Optimization - Summary

### Chapter Overview

Advanced policy optimization methods address stability and sample efficiency challenges in policy learning. This chapter would cover TRPO, PPO, SAC, and population-based approaches.

Key topics include:

- Trust Region Policy Optimization (TRPO)
- Proximal Policy Optimization (PPO)
- Soft Actor-Critic (SAC) and entropy regularization
- Population-based training methods
- Evolutionary strategies for RL



## Part IV

# Advanced Topics



---

This part explores advanced reinforcement learning topics that extend beyond the basic framework. We examine multi-agent systems, hierarchical approaches, model-based methods, and exploration strategies.

These advanced topics require sophisticated mathematical tools and careful analysis of convergence properties, stability, and sample complexity.

---

## Chapter 12

# Multi-Agent Reinforcement Learning - Summary

### Chapter Overview

Multi-agent reinforcement learning addresses scenarios where multiple learning agents interact. This chapter would cover game theory foundations, coordination mechanisms, and emergent behavior.





## Chapter 13

# Hierarchical Reinforcement Learning - Summary

### Chapter Overview

Hierarchical reinforcement learning decomposes complex tasks into simpler subtasks. This chapter would cover options framework, goal-conditioned RL, and skill discovery.



## Chapter 14

# Model-Based Reinforcement Learning - Summary

### Chapter Overview

Model-based reinforcement learning learns environment models to improve sample efficiency. This chapter would cover Dyna-Q, world models, and planning algorithms.



## Chapter 15

# Exploration and Exploitation - Summary

### Chapter Overview

The exploration-exploitation tradeoff is fundamental to reinforcement learning. This chapter would cover bandit algorithms, curiosity-driven exploration, and information-theoretic approaches.



## Part V

# Applications and Frontiers





---

This final part focuses on practical deployment and emerging research directions. We examine transfer learning, real-world applications, and future research frontiers.

---

## Chapter 16

# Transfer Learning and Meta-Learning - Summary

### Chapter Overview

Transfer learning and meta-learning enable rapid adaptation to new tasks. This chapter would cover domain adaptation, MAML, and few-shot learning approaches.



## Chapter 17

# Real-World Applications and Deployment - Summary

### Chapter Overview

Real-world deployment requires careful consideration of safety, robustness, and engineering constraints. This chapter would cover applications in robotics, finance, and autonomous systems.



## Chapter 18

# Future Directions and Research Frontiers - Summary

### Chapter Overview

The field continues to evolve rapidly with new theoretical insights and applications. This chapter would explore emerging paradigms and research directions.





# Conclusion

This enhanced textbook provides a comprehensive foundation in reinforcement learning for engineer-mathematicians. The combination of mathematical rigor, practical implementation guidance, and interactive learning materials creates a unique resource for both students and practitioners.

The first eight chapters provide a solid foundation covering mathematical prerequisites through neural network-based reinforcement learning. The remaining chapters (9-18) outline the complete scope of modern reinforcement learning, from policy methods through cutting-edge research frontiers.

## Next Steps

To continue your reinforcement learning journey:

1. Complete the interactive notebooks for hands-on experience
2. Implement the algorithms in your domain of interest
3. Explore the research literature for specific applications
4. Contribute to the open-source community

## Resources

- **Repository:** <https://github.com/adiel2012/reinforcement-learning>
- **Interactive Notebooks:** Available in Google Colab
- **Community:** Join discussions and contribute improvements

The field of reinforcement learning continues to evolve rapidly. This textbook provides the mathematical foundations and practical skills needed to participate in this exciting journey from theory to application.